

Bayesian Networks Classifiers

Gladys Castillo
University of Aveiro

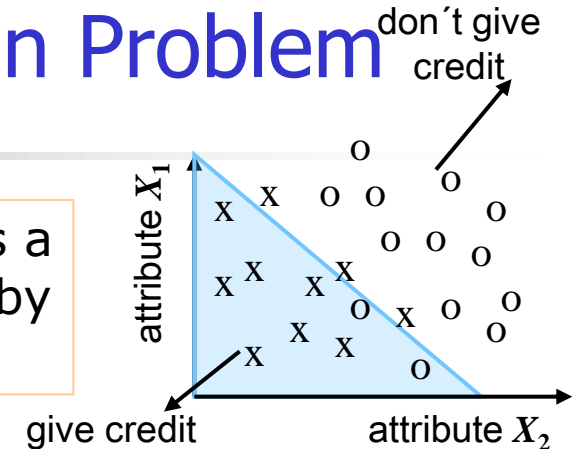


Bayesian Networks Classifiers

Part I – Naive Bayes

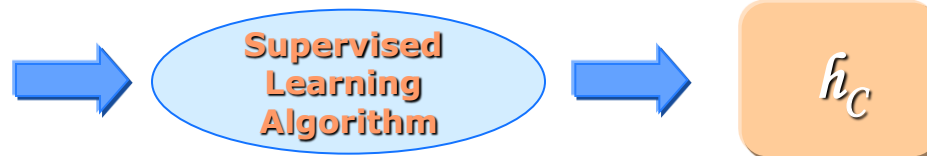
The Supervised Classification Problem

A **classifier** is a function $f: \Omega_{\mathbf{X}} \rightarrow \Omega_C$ that assigns a class label $c \in \Omega_C = \{c_1, \dots, c_m\}$ to objects described by a set of attributes $\mathbf{X} = \{X_1, X_2, \dots, X_n\} \in \Omega_{\mathbf{X}}$



Learning Phase: **Given:** a dataset \mathcal{D} with N labeled examples of $\langle \mathbf{X}, C \rangle$
Build: a **classifier**, a **hypothesis** $h_C: \Omega_{\mathbf{X}} \rightarrow \Omega_C$ that can correctly predict the class labels of new objects

\mathcal{D}	X_1	...	X_n	C
$\langle \mathbf{x}^{(1)}, c^{(1)} \rangle$	$x_1^{(1)}$		$x_n^{(1)}$	$c^{(1)}$
$\langle \mathbf{x}^{(2)}, c^{(2)} \rangle$	$x_1^{(2)}$		$x_n^{(2)}$	$c^{(2)}$
...
$\langle \mathbf{x}^{(N)}, c^{(N)} \rangle$	$x_1^{(N)}$		$x_n^{(N)}$	$c^{(N)}$

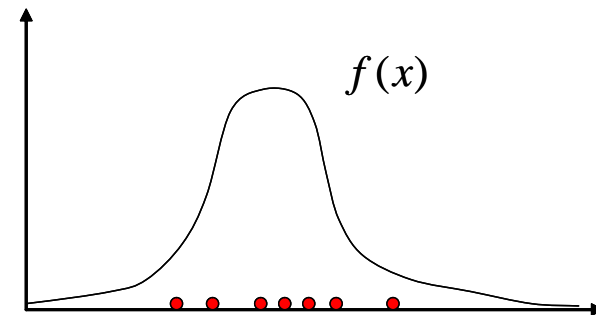


Classification Phase: the class attached to $\mathbf{x}^{(N+1)}$ is $c^{(N+1)} = h_C(\mathbf{x}^{(N+1)}) \in \Omega_C$



Statistical Classifiers

- Treat the attributes $X = \{X_1, X_2, \dots, X_n\}$ and the class C as random variables
 - A random variable is defined by its *probability density function*
- Give the probability $P(c_j | \mathbf{x})$ that \mathbf{x} belongs to a particular class rather than a simple classification



Probability density function of a random variable and few observations

\Rightarrow instead of having the map $\mathbf{X} \rightarrow \mathbf{C}$, we have $\mathbf{X} \rightarrow P(\mathbf{C} | \mathbf{X})$

The class c^* attached to an example is the class with bigger $P(c_j | \mathbf{x})$

$$c^* = h_C(\mathbf{x}) = \arg \max_{j=1 \dots m} P(c_j | \mathbf{x})$$

Bayesian Classifiers

Bayesian because the class c^* attached to an example x is determined by the **Bayes' Theorem**

$$P(X, C) = P(X | C) \cdot P(C)$$

$$P(X, C) = P(C | X) \cdot P(X)$$



$$P(C | X) = \frac{P(C)P(X | C)}{P(X)}$$

Bayes theorem is the main tool in Bayesian inference

We can combine the prior distribution and the likelihood of the observed data in order to derive the posterior distribution.

Bayes Theorem

Example

$$P(C | X) = \frac{P(C)P(X | C)}{P(X)}$$

posterior \propto prior \times likelihood

Before observing the data, our prior beliefs can be expressed in a prior probability distribution that represents the knowledge we have about the unknown features.

After observing the data our revised beliefs are captured by a posterior distribution over the unknown features.

Given:

- ✓ A doctor knows that meningitis causes stiff neck 50% of the time ← likelihood
- ✓ Prior probability of any patient having meningitis is 1/50,000 ← prior
- ✓ Prior probability of any patient having stiff neck is 1/20 ← prior

If a patient has stiff neck,

what's the probability he/she has meningitis?

$$P(M | S) = \frac{P(S | M)P(M)}{P(S)} = \frac{0.5 \times 1/50000}{1/20} = 0.0002$$

Bayesian Classifier

Classify \mathbf{x} to the class which has **bigger posterior probability**

$$c^* = h_C(\mathbf{x}) = \arg \max_{j=1 \dots m} P(c_j | \mathbf{x})$$

Maximum a posteriori classification

How to determine $P(c_j | \mathbf{x})$ for each class c_j ?

Bayes Theorem

$$P(c_j | \mathbf{x}) = \frac{P(c_j)P(\mathbf{x} | c_j)}{P(\mathbf{x})}$$

*$P(\mathbf{x})$ can be ignored because is just a normalizing constant that ensures the posterior adds up to 1; it can be computed by summing up the numerator over all possible values of C , i.e.,
 $P(\mathbf{x}) = \sum_j P(c_j) P(\mathbf{x} | c_j)$*

Bayesian Classification Rule

$$c^* = h_{Bayes}(\mathbf{x}) = \arg \max_{j=1 \dots m} P(c_j)P(\mathbf{x} | c_j)$$

Naïve Bayes (NB) Classifier

- ▶ **"Bayes"** because the class c^* attached to an example \mathbf{x} is determined by the Bayes' Theorem

$$c^* = h_{Bayes}(\mathbf{x}) = \arg \max_{j=1 \dots m} P(c_j) P(\mathbf{x} | c_j)$$

when the attribute space is high dimensional direct estimation is hard unless we introduce some assumptions



- ▶ **"Naïve"** because of its very naïve **independence assumption**:

all the attributes are conditionally independent given the class



$$P(\mathbf{x} | c_j) = \prod_{i=1}^n P(X_i = x_i | c_j)$$

$P(\mathbf{x} | c_j)$ can be decomposed into a product of n terms, one term for each attribute

**NB
Classification
Rule**

$$c^* = h_{NB}(\mathbf{x}) = \arg \max_{j=1 \dots m} P(c_j) \prod_{i=1}^n P(X_i = x_i | c_j)$$

Naïve Bayes (NB)

Learning Phase

(Statistical Parameter Estimation)

Given a training dataset \mathcal{D} of N labeled examples (assuming complete data)

1. Estimate $P(c_j)$ for each class c_j

$$\hat{P}(c_j) = \frac{N_j}{N}$$

N_j - the number of examples of the class c_j

2. Estimate $P(X_i = x_k | c_j)$ for each value x_k of the attribute X_i and for each class c_j

■ X_i discrete

$$\hat{P}(X_i = x_k | c_j) = \frac{N_{ijk}}{N_j}$$

N_{ijk} - number of examples of the class c_j having the value x_k for the attribute X_i

■ X_i continuous

Two options {

- The attribute is **discretized** and then treats as a discrete attribute
- A **Normal distribution** is usually assumed

$$P(X_i = x_k | c_j) = g(x_k; \mu_{ij}, \sigma_{ij})$$

onde

$$g(x; \mu, \sigma) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

The mean μ_{ij} and the standard deviation σ_{ij} are estimated from \mathcal{D}

Continuous Attributes

Normal or Gaussian Distribution

Estimate $P(X_i = x_k | c_j)$ for a value of the attribute X_i and for each class c_j

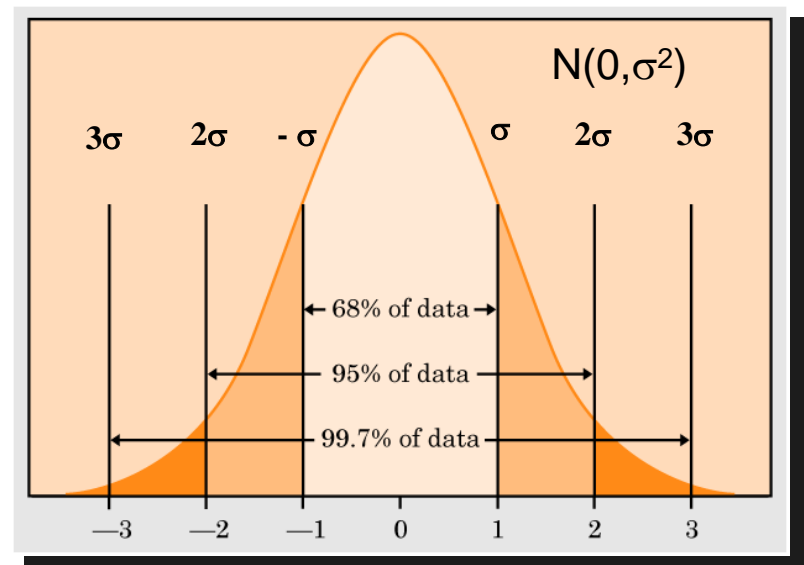
- For real attributes a Normal distribution is usually assumed

$$P(X_i = x_k | c_j) = g(x_k; \mu_{ij}, \sigma_{ij}) \Rightarrow g(x; \mu, \sigma) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

$X_i | c_j \sim N(\mu_{ij}, \sigma_{ij}^2)$ - the mean μ_{ij} and the standard deviation σ_{ij} are estimated from \mathcal{D}

For a variable $X \sim N(74, 36)$, the probability of observing the value 66 is given by:

$$f(x) = g(66; 74, 6) = 0.0273$$



$f(x)$ is symmetrical around its mean

Naïve Bayes

Probability Estimates

Example from John & Langley (1995)

Training dataset \mathcal{D}

Class	X_1	X_2
+	a	1.0
+	b	1.2
+	a	3.0
-	b	4.4
-	b	4.5

Binary Classification Problem

- Two classes: **+** (positive) , **-** (negative)
- Two attributes:
 - X_1 – discrete which takes values a e b
 - X_2 – continuous

1. Estimate $P(c_j)$ for each class c_j

$$\hat{P}(C = +) = \frac{3}{5}$$

$$\hat{P}(C = -) = \frac{2}{5}$$

2. Estimate $P(X_i = x_k / c_j)$ for each value of X_i and each class c_j

X_1 discrete

$$\hat{P}(X_i = a | C = +) = \frac{2}{3}$$

$$\hat{P}(X_i = b | C = +) = \frac{1}{3}$$

$$\hat{P}(X_i = a | C = -) = \frac{0}{2}$$

$$\hat{P}(X_i = b | C = -) = \frac{2}{2}$$

X_2 continuous

$$P(X_2 = x | C = +) = g(x; 1.73, 1.10)$$

$$\mu_{2+} = 1.73, \sigma_{2+} = 1.10$$

$$P(X_2 = x | C = -) = g(x; 4.45, 0.07)$$

$$\mu_{2-} = 4.45, \sigma_{2-}^2 = 0.07$$

The Balance-scale Problem

The dataset was generated to model psychological experimental results

- **Each example has 4 numerical attributes:**

- the left weight (Left_W)
- the left distance (Left_D)
- the right weight (Right_W)
- right distance (Right_D)

- ▶ **Each example is classified into 3 classes:**
the balance-scale:

- ▶ tip to the right (Right)
- ▶ tip to the left (Left)
- ▶ is balanced (Balanced)

- ▶ **3 target rules:**

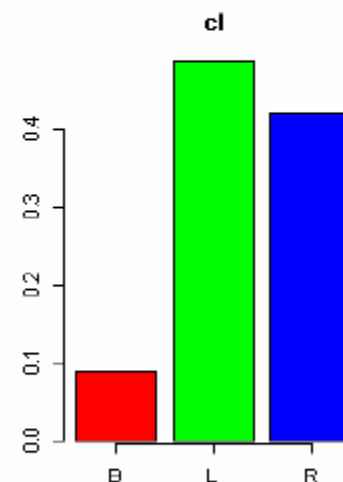
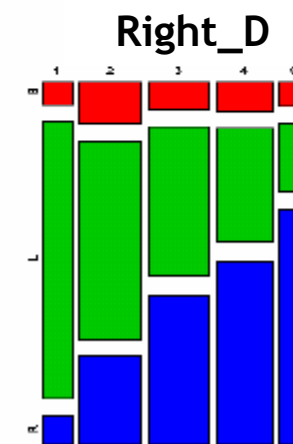
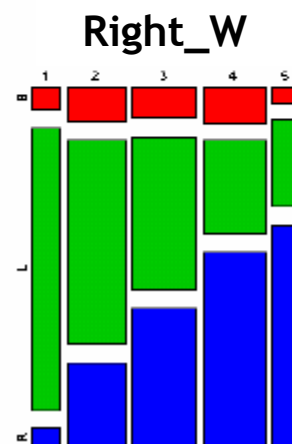
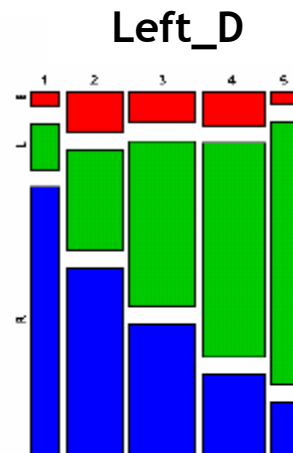
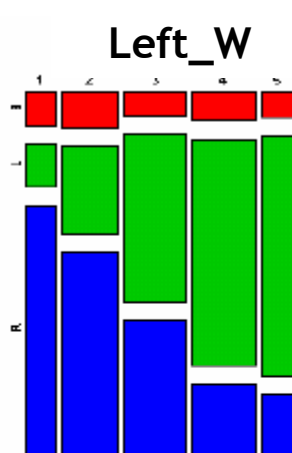
- ▶ If $LD \times LW > RD \times RW \Rightarrow$ tip to the left
- ▶ If $LD \times LW < RD \times RW \Rightarrow$ tip to the right
- ▶ If $LD \times LW = RD \times RW \Rightarrow$ it is balanced



The Balance-scale Problem

Discretization is applied: each attribute is mapped to 5 intervals

Balance-Scale DataSet				
Left_W	Left_D	Right_W	Right_D	Class
1	5	4	2	Right
2	5	3	2	Left
3	4	6	2	Balanced
...



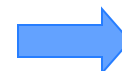
The Balance-scale Problem

Learning Phase

565 examples

Balance-Scale DataSet				
Left_W	Left_D	Right_W	Right_D	Class
1	5	4	2	Right
2	5	3	2	Left
3	4	6	2	Balanced
...

Build



Contingency
tables

Assuming complete data, the computation of all the required estimates requires a simple scan through the data, an operation of time complexity $O(N \times n)$, where N is the number of training examples and n is the number of attributes.

Classes Counters			
Left	Balanced	Right	Total
260	45	260	565

Contingency Tables

Attribute: Left_W					
Class	I1	I2	I3	I4	I5
Left	14	42	61	71	72
Balanced	10	8	8	10	9
Right	86	66	49	34	25

Attribute: Left_D					
Class	I1	I2	I3	I4	I5
Left	16	38	59	70	77
Balanced	8	10	9	10	8
Right	90	57	49	37	27

Attribute: Right_W					
Class	I1	I2	I3	I4	I5
Left	87	63	49	33	28
Balanced	8	10	10	9	8
Right	16	37	58	70	79

Attribute: Right_D					
Class	I1	I2	I3	I4	I5
Left	91	65	44	35	25
Balanced	8	10	8	10	9
Right	17	37	57	67	82

The Balance-scale Problem

Classification Phase

$$c^* = h_{NB}(\mathbf{x}) = \arg \max_{j=1 \dots m} P(c_j) \prod_{i=1}^n P(X_i = x_i | c_j)$$

How NB classifies this example?

Left_W	Left_D	Right_W	Right_D	Class
1	5	4	2	?

The class counters and contingency tables are used to compute the posterior probabilities for each class

We need to estimate the posterior probabilities $P(c_j | \mathbf{x})$ for each class

$$P(c_j | \mathbf{x}) = P(c_j) \times P(\text{Left_W}=1 / c_j) \times P(\text{Left_D}=5 / c_j) \\ \times P(\text{Right_W}=4 / c_j) \times P(\text{Right_D}=2 / c_j), \\ c_j \in \{\text{Left, Balanced, Right}\}$$

The class for this example is the class which has bigger posterior probability

$P(\text{Left} \mathbf{x})$	$P(\text{Balanced} \mathbf{x})$	$P(\text{Right} \mathbf{x})$
0.277796	0.135227	0.586978

max

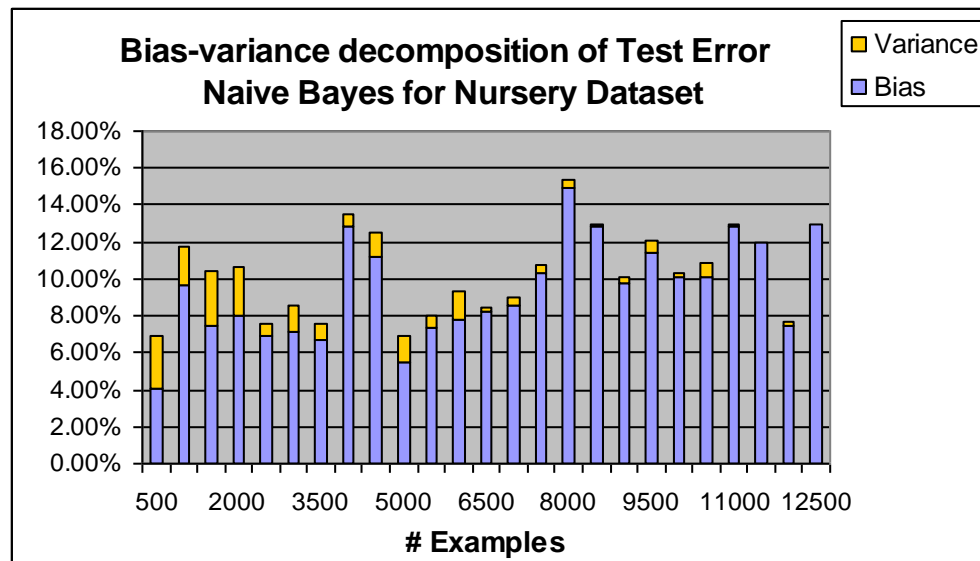
Class = *Right*

Naïve Bayes Performance

- ▶ NB is one of more simple and effective classifiers
- ▶ NB has a very strong unrealistic **independence assumption**:

all the attributes are conditionally independent given the value of class

- ▶ in practice: independence assumption is violated \Rightarrow **HIGH BIAS**
 - ▶ it can lead to poor classification
- ▶ However, **NB is efficient** due to its high variance management
 - ▶ less parameters \Rightarrow **LOW VARIANCE**



Improving Naïve Bayes

- ▶ reducing the bias of **the parameter estimates**

- ▶ by improving the probability estimates computed from data

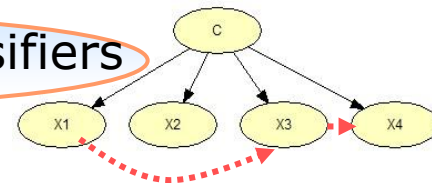
$$\hat{P}(c_j | \mathbf{x}) = \prod_{i=1}^n \hat{P}(X_i = x_i | c_j) \hat{P}(c_j)$$

Relevant
works:

- ▶ Web and Pazzani (1998) - "Adjusted probability naive Bayesian induction" in LNCS v 1502
- ▶ J. Gama (2001, 2003) - "Iterative Bayes", in Theoretical Computer Science, v. 292

- ▶ reducing the bias resulting from **the modeling error**

- ▶ by relaxing the attribute independence assumption
- ▶ one natural extension: **Bayesian Network Classifiers**



Relevant
works:

- ▶ Friedman, Geiger and Goldszmidt (1998) "Bayesian Network Classifiers" in Machine Learning, 29
- ▶ Pazzani (1995) - "Searching for attribute dependencies in Bayesian Network Classifiers" in Proc. of the 5th Workshop of Artificial Intelligence and Statistics
- ▶ Keogh and Pazzani (1999) - "Learning augmented Bayesian classifiers...", in Theoretical Computer Science, v. 292



Bayesian Networks Classifiers

Part II - Introduction to Bayesian Networks

Reasoning under Uncertainty

Probabilistic Approach

- A problem domain is modeled by a set of random variables X_1, X_2, \dots, X_n
- Knowledge about the problem domain is represented by a joint probability distribution $P(X_1, X_2, \dots, X_n)$

Example: ALARM network (Pearl 1988)

- Story: In LA, burglary and earthquake are not uncommon. They both can cause alarm. In case of alarm, two neighbors John and Mary may call
- Problem: Estimate the probability of a burglary based who has or has not called
- Variables: Burglary (B), Earthquake (E), Alarm (A), JohnCalls (J), MaryCalls (M)
- Knowledge required by the probabilistic approach in order to solve this problem: $P(B, E, A, J, M)$

Inference with Joint Probability Distribution



- To specify $P(X_1, X_2, \dots, X_n)$
 - we need at least $2^n - 1$ numbers probability



- exponential model size
- knowledge acquisition difficult (complex, unnatural)
- exponential storage and inference

Solution:



- by exploiting **conditional independence** we can obtain an appropriate factorization of the **joint probability distribution**
 \Rightarrow the number of parameters could be substantially reduced
- A good factorization is provided with **Bayesian Networks**

Bayesian Networks

Bayesian Networks (BNs) graphically represent the joint probability distribution of a set \mathbf{X} of random variables in a problem domain

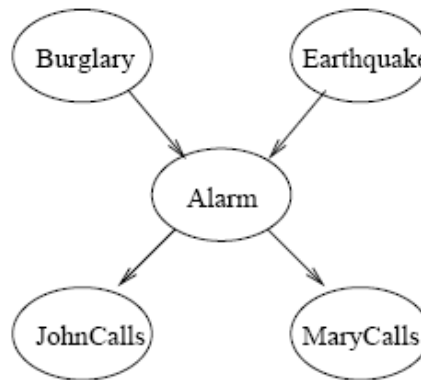
A BN = (S, Θ_S) consist in two components:

Graph theory

Qualitative part

the structure S - a Directed Acyclic Graph (DAG)

- nodes – random variables
- arcs - direct dependence between variables



Probability theory

Quantitative part

the set of parameters Θ_S
= set of conditional probability distributions (CPDs)

- discrete variables:
 Θ_S - multinomial (CPDs are CPTs)
- continuous variables:
 Θ_S - Normal ou Gaussian

Markov

condition : *each node is independent of its non descendants given its parents in S*

⇒ The joint distribution is factorized

$$P(X_1, \dots, X_n) = \prod_i P(X_i \mid pa(X_i))$$

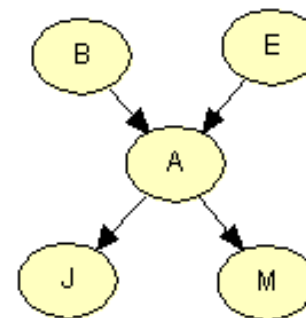
Example: Alarm Network (Pearl)

Bayesian Network

$X = \{ \mathbf{B} - \text{Burglary}, \mathbf{E} - \text{Earthquake}, \mathbf{A} - \text{Alarm}, \mathbf{J} - \text{JohnCalls}, \mathbf{M} - \text{MaryCalls} \}$

The Bayesian Network has two components:

- the structure - a Directed Acyclic Graph



- the parameters - a set of conditional probability tables (CPTs)

$\Rightarrow pa(\mathbf{B}) = \{ \}, \quad pa(\mathbf{E}) = \{ \}, \quad pa(\mathbf{A}) = \{ \mathbf{B}, \mathbf{E} \}, \quad pa(\mathbf{J}) = \{ \mathbf{A} \}, \quad pa(\mathbf{M}) = \{ \mathbf{A} \}$

$P(\mathbf{B})$

B	P(B)
Y	.01
N	.99

$P(\mathbf{E})$

E	P(E)
Y	.02
N	.98

$P(\mathbf{A} \mid \mathbf{B}, \mathbf{E})$

A	B	E	P(A B, E)
Y	Y	Y	.95
N	Y	Y	.05
Y	Y	N	.94
N	Y	N	.06
Y	N	Y	.29
N	N	Y	.71
Y	N	N	.001
N	N	N	.999

$P(\mathbf{J} \mid \mathbf{A})$

J	A	P(J A)
Y	Y	.7
N	Y	.3
Y	N	.01
N	N	.99

$P(\mathbf{M} \mid \mathbf{A})$

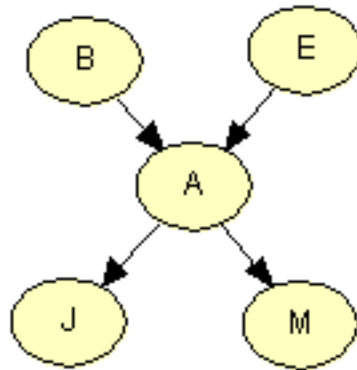
M	A	P(M A)
Y	Y	.9
N	Y	.1
Y	N	.05
N	N	.95

Model size reduced from
31 to $1+1+4+2+2=10$

Example: Alarm Network (Pearl)

Factored Joint Probability Distribution

$X = \{B \text{ (Burglary)}, (E) \text{ Earthquake}, (A) \text{ Alarm}, (J) \text{ JohnCalls}, (M) \text{ MaryCalls}\}$

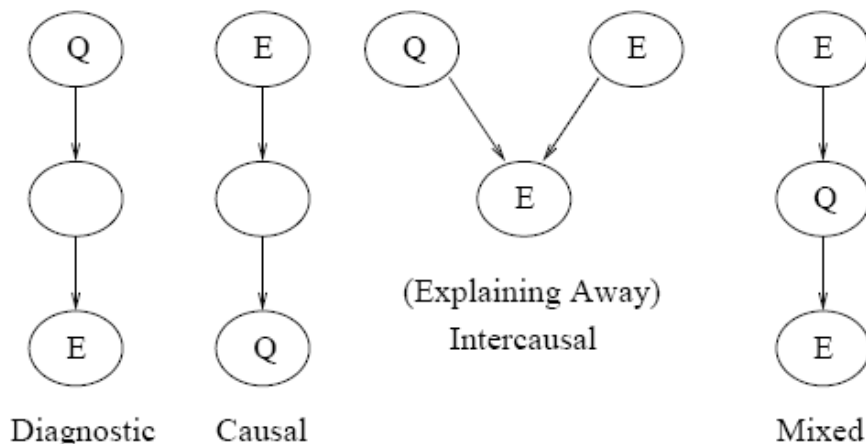


$$P(B, E, A, J, M) = P(B) P(E) P(A|B, E) P(J|A) P(M|A)$$

Inference in Bayesian Networks

Compute the posterior probability distribution for a set of **query variables**, given values for some **evidence variables**

$$P(Q \mid E) = ?$$



Diagnostic inferences: from effect to causes

$P(\text{Burglary} \mid \text{JohnCalls})$

Causal Inferences: from causes to effects

$P(\text{JohnCalls} \mid \text{Burglary})$

$P(\text{MaryCalls} \mid \text{Burglary})$

Intercausal Inferences:

between causes of a common effect

$P(\text{Burglary} \mid \text{Alarm})$

$P(\text{Burglary} \mid \text{Alarm Earthquake})$

Mixed Inference:

combining two or more of above.

$P(\text{Alarm} \mid \text{JohnCalls Earthquake})$

$P(\text{Burglary} \mid \text{JohnCalls Earthquake})$



Bayesian Network Resources

- **Repository:** www.cs.huji.ac.il/labs/compbio/Repository/
- **Softwares:**
for a updated list visit <http://www.kdnuggets.com/software/bayesian.html>
 - GeNIe: genie.sis.pitt.edu
 - Hugin: www.hugin.com
 - Analytica: www.lumina.com
 - JavaBayes: [www.cs.cmu.edu/ javabayes/Home/](http://www.cs.cmu.edu/javabayes/Home/)
 - Bayesware: www.bayesware.com
- **Others**
 - Bayesian Belief Nets by Russell Greiner
<http://webdocs.cs.ualberta.ca/~greiner/bn.html>
 - List of Bayesian Network software by Kevin Murphy
<http://www.cs.ubc.ca/~murphyk/Software/bnsoft.html>



Bayesian Networks: Summary

Bayesian Networks: an **efficient** and **effective** representation of the joint probability distribution of a set of random variables

- Efficient:
 - Local models
 - Independence (d-separation)
- Effective:

Algorithms take advantage of structure to

 - Compute posterior probabilities
 - Compute most probable instantiation
 - Decision making
- But there is more: statistical induction ▲ LEARNING

adapted from © 1998, Nir Friedman, U.C. Berkeley, and Moises Goldszmidt, SRI International. All rights reserved.



Bayesian Network Classifiers

Part III - Learning Bayesian Networks

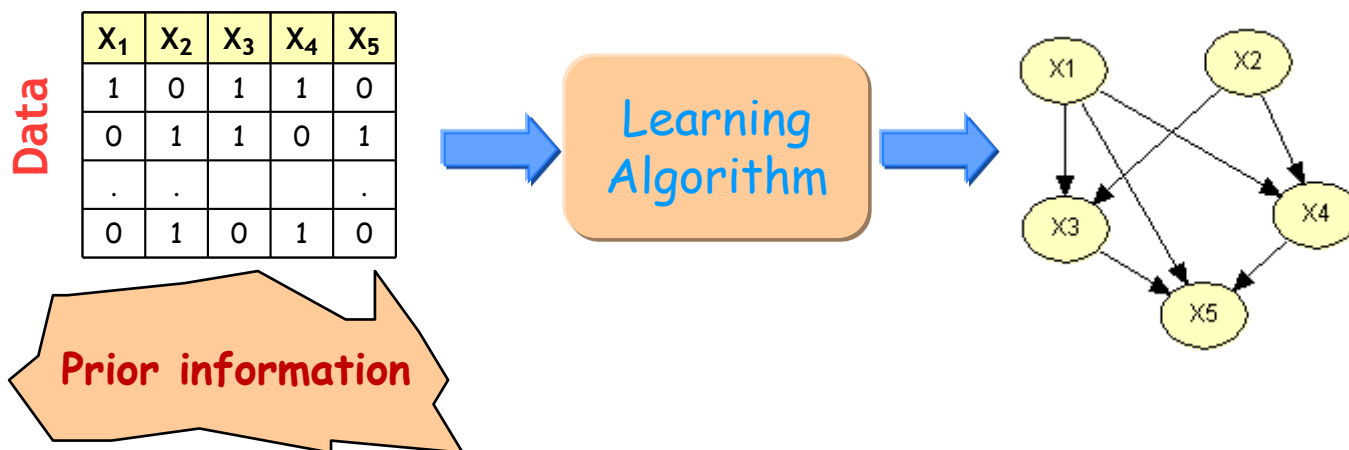
Learning Bayesian Networks

Let $\mathbf{X} = \{X_1, X_2, \dots, X_n\}$ be a set of random variables for a domain under study

Given:

- ✓ a training dataset $\mathcal{D} = \{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(N)}\}$ of i.i.d examples of \mathbf{X}
- ✓ some prior information ξ (background knowledge)
a BN or fragment of it, prior probabilities, etc.

Induce: a Bayesian Network $\text{BN} = (S, \Theta_\xi)$ that best matches \mathcal{D}





Learning Bayesian Networks

How many Learning Problems?

We can distinguish a variety of learning problems, depending on whether the **structure** is *known* or *unknown*, the **data** is *complete* or *incomplete*, and there are *hidden variables* or not.

	Known Structure	Unknown Structure
Complete Data	Statistical parametric estimation (closed-form eq.)	Discrete optimization over structures (discrete search)
Incomplete Data	Parametric optimization (EM, gradient descent...)	Combined (Structural EM, mixture models...)

Adapted from [Nir Friedman and Moises Goldszmidt slides](#)

Learning Problem

Complete Data, Known Structure


	Known Structure	Unknown Structure
Complete	Statistical parametric estimation (closed-form eq.)	Discrete optimization over structures (discrete search)
Incomplete	Parametric optimization (EM, gradient descent...)	Combined (Structural EM, mixture models...)



Parameter Learning

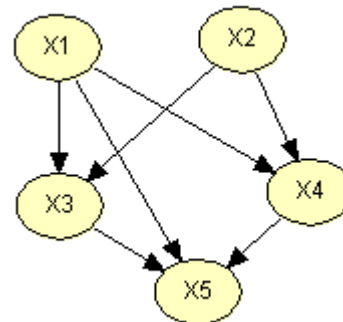
Parameter Learning

Complete Data

	Known Structure	Unknown Structure
Complete data		
Incomplete data		

Given:

- A Bayesian Network Structure



- Complete Data (no missing values)

X ₁	X ₂	X ₃	X ₄	X ₅
1	0	1	1	0
0	1	1	0	1
1	1	0	1	1
.	.			.
0	1	0	1	0

Estimate:

- conditional probabilities:

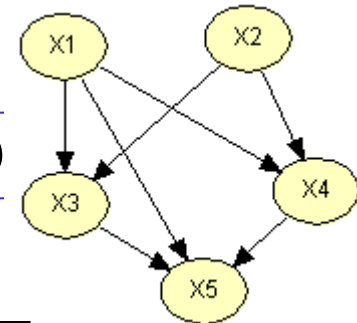
$$P(X_1), P(X_2), P(X_3 | X_1, X_2), P(X_4 | X_1, X_2), P(X_5 | X_1, X_3, X_4)$$

Parameter Learning

Complete Data

Conditional Probabilities Tables:

$$P(X_1), P(X_2), P(X_3 | X_1, X_2), P(X_4 | X_1, X_2), P(X_5 | X_1, X_3, X_4)$$



P(X ₁)	
X ₁ =1	X ₁ =0
?	?

P(X ₂)	
X ₂ =1	X ₂ =0
?	?

pa(X ₃)		P(X ₃ pa(X ₃))	
X ₁	X ₂	X ₃ =1	X ₃ =0
1	1	?	?
1	0	?	?
0	1	?	?
0	0	?	?

pa(X ₄)		P(X ₄ pa(X ₄))	
X ₁		X ₄ =1	X ₄ =0
1		?	?
0		?	?

pa(X ₅)			P(X ₅ pa(X ₅))	
X ₁	X ₃	X ₄	X ₅ =1	X ₅ =0
1	1	1	?	?
1	1	0	?	?
1	0	1	?	?
1	0	0	?	?
0	1	1	?	?
0	1	0	?	?
0	0	1	?	?
0	0	0	?	?

For each parent configuration we have an independent estimation problem for each local binomial model

For binary variables is enough to fill only one probability

Parameter Learning

Complete Data

$\mathbf{X} = \{X_1, X_2, \dots, X_n\}$ - set of random discrete variables

Statistical Parameter Estimation Problem

How to obtain from data an estimate of each conditional probability?

$$\hat{\theta}_{ijk} = \hat{P}(X_i = x_k \mid Pa_i = pa_j) = ?$$

- Estimation relies on **sufficient statistics**

- count the number of cases in \mathcal{D} such that $X_i = x_k$ and $pa(X_i) = pa_j$

$$N_{ijk} \equiv N(X_i = x_k \mid pa(X_i) = pa_j)$$

- count the number of cases in \mathcal{D} such that $pa(X_i) = pa_j$

$$N_{ij} \equiv N(pa(X_i) = pa_j)$$

Parameter Learning

Complete Data

$\mathbf{X} = \{X_1, X_2, \dots, X_n\}$ - set of random discrete variables

$$\hat{\theta}_{ijk} = \hat{P}(X_i = x_k \mid Pa_i = pa_j) = ?$$

- There are two main approaches for parameter estimation:

frequentist

vs.

Bayesian



MLE estimator

Bayesian estimator

$$\hat{\theta}_{ijk} = \frac{N_{ijk}}{N_{ij}}$$

$$\hat{\theta}_{ijk} = \frac{N_{ijk} + \alpha_{ijk}}{N_{ij} + \alpha_{ij}}$$

$$N_{ij} = \sum_k N_{ijk}; \alpha_{ij} = \sum_k \alpha_{ijk}$$

Theoretically they can be thought of as “imaginary” counts from our past experience (priors on counters). In practice they can be thought of as initial counters

Parameter Learning

Complete Data. Example

X_1	X_2	X_3	X_4	X_5
1	0	1	1	1
1	1	1	0	0
0	1	1	1	0
1	1	0	0	1
0	0	1	1	0
0	1	1	0	0
1	0	0	1	1
0	1	1	0	1
1	1	0	1	0
0	1	1	0	0
0	1	0	1	0
1	1	0	0	0
0	1	0	0	1

		pa(X_3)		P(X_3 pa(X_3))	
		X_1	X_2	$X_3=1$	$X_3=0$
pa ₁	→	1	1	?	?
pa ₂	→	1	0	?	?
pa ₃	→	0	1	?	?
pa ₄	→	0	0	?	?

MLE for $P(X_3 = 1 \mid X_1 = 0, X_2 = 1)$

$$\hat{\theta}_{331} = \frac{4}{6} = 0.66$$

MLE for $P(X_3 = 0 \mid X_1 = 0, X_2 = 1)$

$$\hat{\theta}_{330} = 1 - 0.66 = 0.34$$

Learning Problem

Incomplete Data, Known Structure

	Known Structure	Unknown Structure
Complete	Statistical parametric estimation (closed-form eq.)	Discrete optimization over structures (discrete search)
Incomplete	Parametric optimization (EM, gradient descent...)	Combined (Structural EM, mixture models...)

- Methods for learning: EM and Gradient Ascent

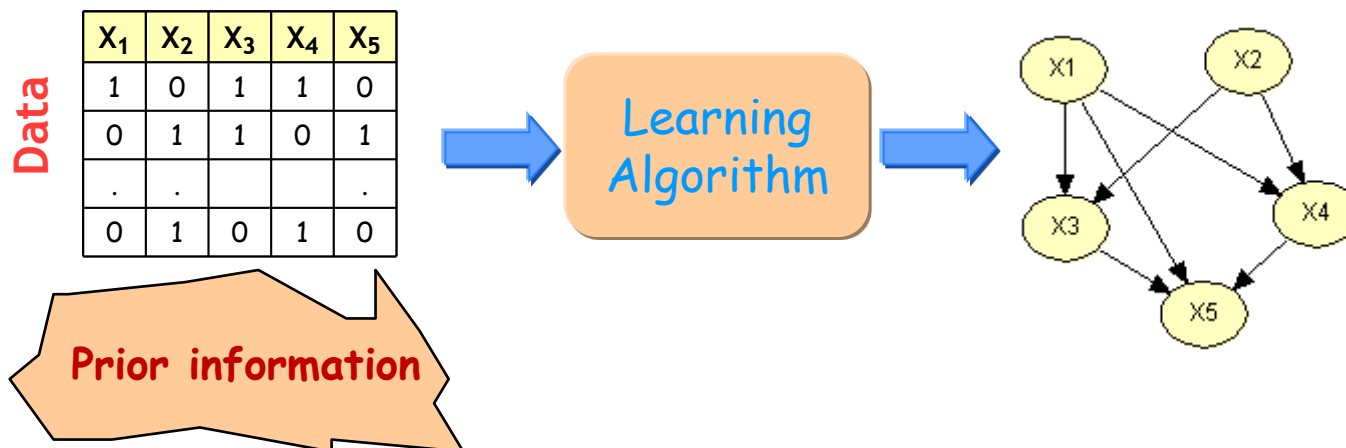
Difficulties:

- Exploration of a complex likelihood/posterior
 - More missing data \Rightarrow many more local maxima
 - Cannot represent posterior \Rightarrow must resort to approximations

Learning Problem

Complete Data, Unknown Structure

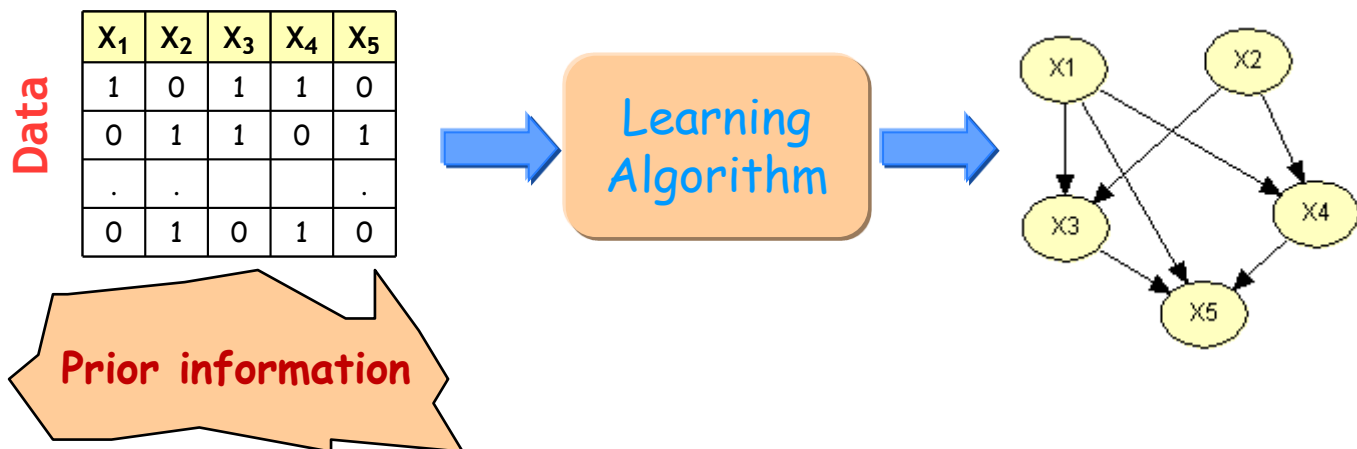
	Known Structure	Unknown Structure
Complete	Statistical parametric estimation (closed-form eq.)	Discrete optimization over structures (discrete search)
Incomplete	Parametric optimization (EM, gradient descent...)	Combined (Structural EM, mixture models...)



Learning Problem

Incomplete Data, Unknown Structure

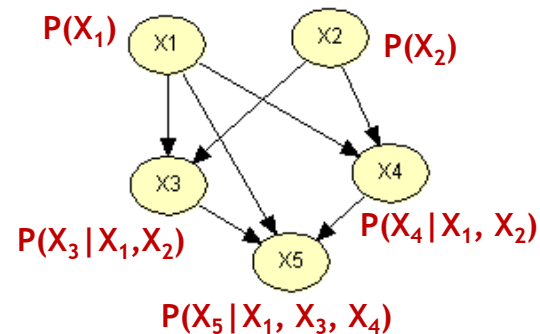
	Known Structure	Unknown Structure
Complete	Statistical parametric estimation (closed-form eq.)	Discrete optimization over structures (discrete search)
Incomplete	Parametric optimization (EM, gradient descent...)	Combined (Structural EM, mixture models...)



Learning Bayesian Networks

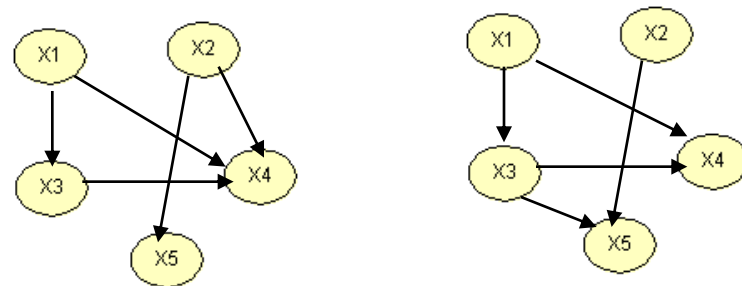
■ Known structure - learn parameters

- **Complete data:**
parameter estimation (ML, MAP)
- **Incomplete data:**
non-linear parametric
optimization (gradient descent, EM)



■ Unknown structure - learn graph and parameters

- **Complete data:**
optimization (search
in space of graphs)
- **Incomplete data:**
structural EM,
mixture models



$$S = \arg \max_S \text{Score}(S)$$



Structure Learning

Complete Data

Statistical Model Selection Problem

Given data, to find the structure that best fits the data

Two approaches:

- Constraint based (dependency analysis & search)
 - Perform tests of conditional independence
 - χ^2 test or mutual information test locally measure the dependency relationships between the variables
 - Search for a network that is consistent with the observed dependencies and independencies
- Score based (scoring & search)
 - Define a score that evaluates how well the (in)dependencies in a structure match the observations
 - Search for a structure that maximizes the score



Structure Learning

Score based Approaches

- Must choose:

a ***score(S)*** to measure how well a model fits the data

Discrete Optimization Problem: Find the structure that maximizes the score in the space \mathcal{S} of possible networks

$$\hat{\mathcal{S}} = \arg \max_{\mathcal{S}} \text{score}(\mathcal{S})$$

NP-hard problem
proved by Chickering
et al. (1994)

- NP-hard optimization problem:

we can solve it by using ***heuristic search algorithms***

- greedy hill climbing, best first search, simulated annealing, etc.



Structure Learning Scores

From a practical point of view, not philosophical we can classify scores into four categories:

1. Log-likelihood-based scores
 - **MLC** - Maximum Likelihood criterion
2. Penalized log-likelihood scores
 - **BIC** - Bayesian Information Criterion (Schwarz, 1978)
 - **MDL** - Minimum Description Length (Rissanen, 1978)
 - **AIC** - Akaike's Information Criterion
3. Bayesian scores: **Bayes**, **BDe**, **K2**
4. Predictive scores
 - Cross-validation (**k-Fold-CV**)
 - Prequential (**Preq**)

Structure Learning

Log-likelihood based Scores

$$\mathcal{L}(M : \mathcal{D}) \equiv P(\mathcal{D} | M)$$

- **MLC score:** derived from the Fisher's likelihood principle:

“a hypothesis is more plausible or likely than another, in the light only of the observed data if it makes those data more probable”

$$Score_{MLC}(S, \mathcal{D}) \equiv l(\hat{\Theta}_{ML} : \mathcal{D}, S) = \sum_{i=1}^n \sum_{j=1}^{q_i} \sum_{k=1}^{r_i} N_{ijk} \log \frac{N_{ijk}}{N_{ij}}$$

- **Penalized log-likelihood scores: BIC/MDL, AIC**

MDL e AIC are derived from information-theoretic arguments

$$Score_{BIC}(S, \mathcal{D}) \equiv Score_{MLC}(S, \mathcal{D}) - \frac{1}{2} \log N \parallel S \parallel$$

*maximize BIC is
equivalent
to minimize MDL*

$$Score_{AIC}(S, \mathcal{D}) \equiv Score_{MLC}(S, \mathcal{D}) - \parallel S \parallel$$

$\parallel S \parallel$ is the dimension of the BN defined as the number of its paramters

Structure Learning

Bayesian Scores

Bayes (BD) score: the log of the relative posterior probability

$$Score_{BD}(S^h, \mathcal{D}) \equiv \log P(S^h) + \log P(\mathcal{D} | S^h)$$

To obtain the Bayesian score we need to assess the **prior distribution** $P(S)$ for each candidate structure and to compute **the marginal likelihood** $P(\mathcal{D} | S)$

Cooper et al (1992) and Heckerman et al.(1995) proved that under some nice assumptions the marginal likelihood can be derived in a closed-form solution and it decomposes into a product of terms, one for each local distribution

$$P(\mathcal{D} | S^h) = \prod_{i=1}^n \prod_{j=1}^{q_i} \frac{\Gamma(\alpha_{ijk})}{\Gamma(\alpha_{ij} + N_{ij})} \prod_{k=1}^{r_i} \frac{\Gamma(\alpha_{ijk} + N_{ijk})}{\Gamma(\alpha_{ijk})}$$

where $\alpha_{ij} = \sum_{k=1}^{r_i} \alpha_{ijk}$, $N_{ij} = \sum_{k=1}^{r_i} N_{ijk}$ and $\Gamma(\cdot)$ is the gamma-function.

↓
Priors on counters = initial counters



Structure Learning

Bayesian Scores

$$Score_{BD}(S^h, \mathcal{D}) \equiv \log P(S^h) + \log P(\mathcal{D} | S^h)$$

To compute the Bayes(BD) score we need to set:

1. Priors over structures

- Assuming uniform priors for all the candidate we obtain the **log marginal likelihood**
(one of the most commonly used scores in learning BNs)

2. Priors on parameters:

different priors leads to special cases of the BD score:

- **K2** (Cooper and Herskovits, 1992): *log marginal likelihood* with the simple *non-informative* prior $\alpha_{ijk} = 1$
- **BDe** (Buntine, 1991; Heckerman et al., 1995):
BD with the additional assumption of *likelihood equivalence*
 - $\alpha_{ijk} = 1/(r_i \cdot q_i)$, r_i – the number of possible values of X_i
 q_i – the number of possible parent configurations of X_i

Structure Learning

Bayesian Scores

Bayes (BD) score:

$$Score_{BD}(S^h, \mathcal{D}) \equiv \log P(S^h) + \log P(\mathcal{D} | S^h)$$

To obtain the Bayesian score we need to assess the **prior distribution** $P(S)$ for each candidate structure and to compute **the marginal likelihood** $P(\mathcal{D} | S)$

Cooper et al.(1992) and Heckerman et al (1995) proved that under some nice assumptions $P(\mathcal{D} | S^h)$ can be derived in a closed-form solution and it decomposes into a product of terms, one for each local distribution

$$P(\mathcal{D} | S^h) = \prod_{i=1}^n \prod_{j=1}^{q_i} \frac{\Gamma(\alpha_{ijk})}{\Gamma(\alpha_{ij} + N_{ij})} \prod_{k=1}^{r_i} \frac{\Gamma(\alpha_{ijk} + N_{ijk})}{\Gamma(\alpha_{ijk})}$$

where $\alpha_{ij} = \sum_{k=1}^{r_i} \alpha_{ijk}$, $N_{ij} = \sum_{k=1}^{r_i} N_{ijk}$ and $\Gamma(\cdot)$ is the gamma-function.

↓
Priors on counters = initial counters

Structure Learning

Predictive Scores

Given: a set $\mathcal{S} = \{S_1, S_2, \dots, S_m\}$ of candidate BN structures

Choose: $S \in \mathcal{S}$ so that the predictive distribution yields the most accurate predictions for future data

Predictive scores require a **loss function** for measuring the predictive accuracy

unsupervised learning (BN is used for general purposes)

The training dataset $\mathcal{D} = \{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(N)}\}$ has N i.i.d examples of \mathbf{X} :

$$\log \text{Loss}(S, \mathcal{D}) = -\sum_{l=1}^N \log P(\mathbf{x}^{(l)} | S, \theta_S)$$

The log-loss of S is based on the joint predictive distribution $P(\mathbf{X})$
 \Rightarrow *generative model*

supervised learning (BN is used as a classifier)

The training dataset $\mathcal{D} = \{<\mathbf{x}^{(1)}, c^{(1)}>, \dots, <\mathbf{x}^{(N)}, c^{(N)}>\}$ has N i.i.d labelled examples

✓ **the zero-one loss** \leftarrow *the most used*

✓ the conditional log-loss

$$C \log \text{Loss}(S, \mathcal{D}) = -\sum_{l=1}^N \log P(c^{(l)} | \mathbf{x}^{(l)}, S, \theta_S)$$

The log-loss of S is based on the conditional predictive distribution $P(\mathbf{X}/C)$
 \Rightarrow *discriminative model*

Structure Learning

k-fold Cross-Validation Score

One natural way to measure the predictive performance is provided by **cross-validation** (Stone -1974)

Algorithm 1 The algorithm for computing the k -fold cross-validation score for Bayesian networks

Require: A Bayesian network structure S , a dataset \mathcal{D} of *i.i.d.* examples of \mathbf{X} , a loss function $\text{lossF}(\text{BN}, \mathcal{D})$, the number of folds k

Ensure: The k-Fold-CV score for the structure S given the data \mathcal{D}

1: Split the dataset \mathcal{D} in k folds

2: $\Theta_S \leftarrow \text{initialize-CPTs}(S)$

3: $\text{BN} \leftarrow (S, \Theta_S)$

4: **for** each fold in \mathcal{D} **do**

5: $\mathcal{D}_{\text{training}} \leftarrow \mathcal{D} \setminus \text{fold \{first: training\}}$

← **training**

6: $\text{learnParameters}(\Theta_S, \mathcal{D}_{\text{training}})$

7: $\mathcal{D}_{\text{test}} \leftarrow \text{fold \{second: testing\}}$

← **testing**

8: $\text{loss}[\text{fold}] \leftarrow \text{lossF}(\text{BN}, \mathcal{D}_{\text{test}})$

9: **end for**

10: **return** $\text{Average}(\text{loss}[\text{fold}])$ {the k-Fold-CV score for S given data \mathcal{D} }

The k-fold-CV score is the average over the k-loss

The **leave-one-out cross-validation score** (LOO-CV) is a particular case when $k = N$

Structure Learning

A. P. Dawid, 1984

Prequential (Cumulative) Score

The prequential score is based on the Dawid's prequential approach

Algorithm 2 The algorithm for computing the prequential score for Bayesian networks

Require: A Bayesian network structure S , a dataset \mathcal{D} of *i.i.d.* examples of \mathbf{X} , a loss function $\text{lossF}(\text{BN}, \mathcal{D})$

Ensure: The *prequential* score Preq for the structure S given the data \mathcal{D}

```
1:  $\Theta_S \leftarrow \text{initialize-CPTs}(S)$ 
2:  $\text{BN} \leftarrow (S, \Theta_S)$ 
3: for each example  $\mathbf{x}$  in  $\mathcal{D}$  do
4:    $\text{cumLoss} \leftarrow \text{cumLoss} + \text{lossF}(\text{BN}, \mathbf{x})$  {first: predict}
5:    $\text{update}(\Theta_S, \mathbf{x})$  {second: update the parameters with new example}
6: end for
7: return  $\text{cumLoss}$  {the prequential score for  $S$  given data  $\mathcal{D}$ }
```

First the current model is used to do prediction. Then the example is used to update the parameters

The prequential score is the cumulative loss



Structure Learning

Scores. Conclusions

- MDL and Bayesian scores are the most popular in learning BNs
Asymptotically: MDL is equivalent to MLC and BDe
- Evaluation of prior distributions over structures and parameters
 - Log-likelihood based scores (frequentist approach): do not require
 - Bayesian scores (bayesian approach to model selection): require
- Model Complexity (number of parameters)
 - Likelihood based scores prefer more complex models
⇒ can **overfit** the data, specially for small training dataset
 - Penalized likelihood scores (BIC/MDL, AIC) are based on the *Occam's Razor*:
“*given two equally predictive theories, choose the simpler*”
⇒ a more optimal *complexity-fitness* trade-off but can **underfit** for small data
- Arranged according their bias toward simplicity:
MDL, Bayes, AIC, Preq, CV, MLC, BDe

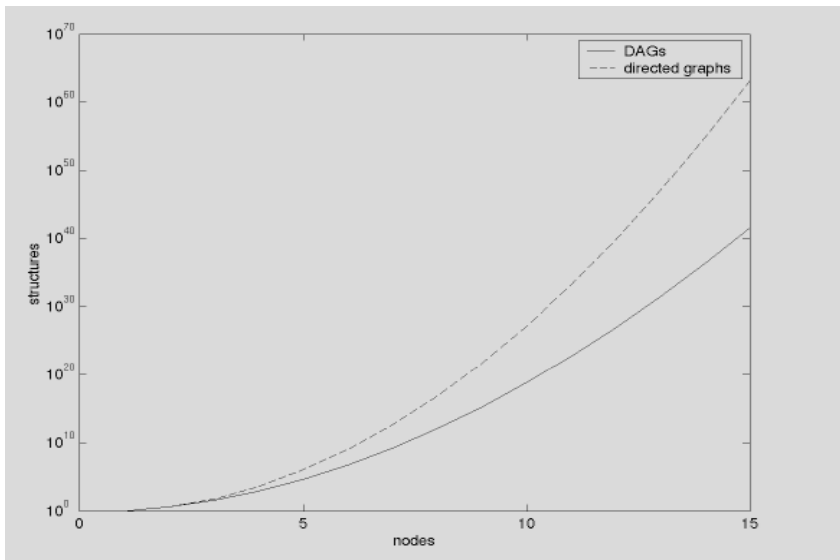
Structure Learning

Search Problem

a discrete optimization problem

Problem Optimization:

How to find the structure that maximizes a scoring function?



Naïve Method: **brute-force**

1. compute the score of every structure
2. pick the one with the highest score

The number of possible structures grows super-exponentially with the number of variables \Rightarrow exhaustive search is infeasible



The problem of finding the best structure is known to be **NP-hard**
 \Rightarrow use **heuristic search algorithms** to traverse the solution space
in searching for optimal/near optimal solutions



Structure Learning

Heuristic Search Algorithms

We need to define:

- Search space:
 - solution space
 - set of operators
- Initial solution in the solution space
- Search strategy
- Objective Function: a scoring function
- Goal State: usually a stopping criterion is used
 - Ex: stop when the new solution cannot improve the current solution

Structure Learning

Search Space

B-space

the simplest formulation of the search space - (Chickering - 2002)

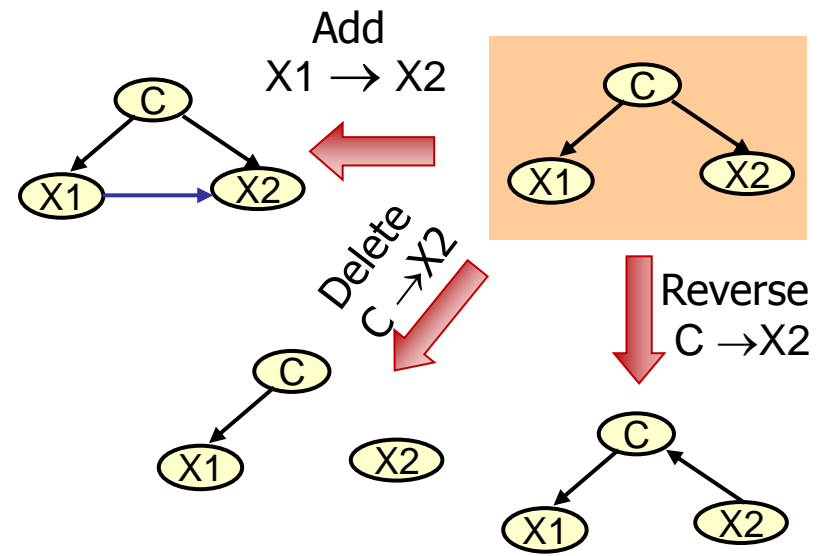
- **Solution Space:** the set of all possible solutions

- individual DAGs

- **Set of Operators:**

used by the search algorithm to transform one state to another

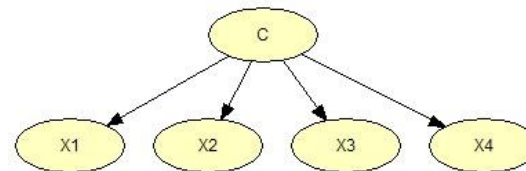
- arc addition: **addArc**
- arc deletion: **deleteArc**
- arc reversion: **reverseArc**



Structure Learning

B-space: Initial Solution

- a DAG with no edges
then iteratively add arcs that most increase the score subject to never introducing a cycle
- a complete DAG
then iteratively delete arcs that most increase the score subject to never introducing a cycle
- Some DAG in the middle of the B-space
ex: a Naive Bayes structure





Structure Learning

Search Strategies

Search strategies define how to organize the search in the search space:

- **deterministics**: all the runs always obtain the same solution
 - Hill-climbing (greedy search) (one of the most used in learning BNs)
 - Tabu search
 - Branch and bound

As a rule, they tend to get stuck in local maximums

- **non-deterministics**:
use randomness for escaping from local maximums
 - simulated annealing
 - Genetic algorithms
 - GRASP - Greedy Randomize Adaptive Search procedure

Structure Learning

Local score based search

Local search methods:

use scores that are decomposed into local scores, one for each node

$$Score(S, \mathcal{D}) = \sum_{i=1}^n Score_{local}(X_i \mid Pa(X_i), N_{X_i \mid Pa(X_i)})$$

- change in score that results from the application of an operator can be computed locally
- the global optimization problem is decomposed into local optimization problems
- Local score metrics: **MLC**, **AIC**, **BIC/MDL**, **BD**, **K2**, **BDe** (all are decomposable)
 - They are based on the **log likelihood** or **log marginal likelihood**
both decompose into a sum of terms, one for each node



Structure Learning

Global score based search

- Global search methods:

use scores that cannot be decomposed into local scores for each node

- Predictive scores - **k-Fold-CV**, **Prequential**, etc.
- the whole network needs to be considered to determine the change in the score that results from the application of an operator



Structure Learning

Greedy Hill Climbing

Algorithm 3 The hill-climbing search algorithm for learning the structure of Bayesian networks

Require: A B-space (S, \mathcal{O}) , where S is the space of possible DAGs, and $\mathcal{O} = \{\text{addArc}, \text{deleteArc}, \text{reverseArc}\}$ is the set of possible operators, an initial structure $S \in S$, a dataset \mathcal{D} of *i.i.d.* examples of a set $\mathbf{X} = \{X_1, X_2, \dots, X_n\}$ of random variables for a domain under study, a scoring function $\text{Score}(S, \mathcal{D})$

Ensure: A Bayesian network structure $S \in S$ with high value of the score

```
1: continue  $\leftarrow$  True
2: while continue do
3:   Compute  $\text{Score}(S, \mathcal{D})$ 
4:   Find best operator  $\text{op}$  such that  $\text{op} = \arg \max_{\text{op} \in \mathcal{O}} \text{Score}(\text{op}(S), \mathcal{D})$ 
5:   if  $\text{op}$  exists  $\wedge \text{Score}(\text{op}(S), \mathcal{D}) > \text{Score}(S, \mathcal{D})$  then
6:      $S \leftarrow \text{op}(S)$  {Apply the operator to the current structure}
7:   else
8:     continue  $\leftarrow$  False
9: end while
10: return  $S$  {a structure with a high score}
```

Init: some initial structure

at each search step:

- ▶ apply the operator that more increases the score

STOP searching:

- ▶ when there is no more improvement of the score
- ▶ when it is no possible to add a new arc



Structure Learning

Problems with hill-climbing

- Local maxima:
 - All one-edge changes reduced the score, but not optimal yet.
- Plateaus:
 - Neighbors have the same score.
- Solutions:
 - Random restart.
 - TABU-search:
 - Keep a list of K most recently visited structures and avoid them.
 - Avoid plateau.
 - Simulated annealing.

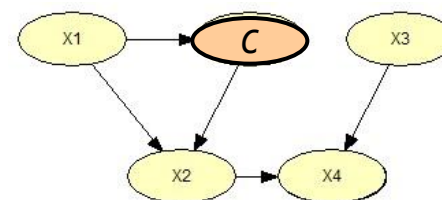


Bayesian Network Classifiers

Part IV - Learning Bayesian Network Classifiers

Bayesian Network Classifiers

- In classification problems the domain variables are partitioned into:
 - the set of attributes, $\mathbf{X} = \{X_1, \dots, X_n\}$
 - the class variable C



- A BN can be used as a classifier that gives the **posterior probability distribution of the class node C** given the values of the attributes \mathbf{X}
 - Given an example \mathbf{x} we can compute the predictive distribution $P(C | \mathbf{x}, S)$ by marginalizing the joint distribution $P(C, \mathbf{x} | S)$

$$P(C | \mathbf{x}, S) = \frac{P(C, \mathbf{x} | S)}{P(\mathbf{x} | S)} \propto P(C, \mathbf{x} | S)$$

Bayesian Network Classifiers

Classification Rule

BNC
Classification
Rule

$$c^* = h_{BNC}(\mathbf{x}) = \arg \max_{j=1 \dots m} P(\mathbf{x}, c_j \mid S, \Theta_s)$$

How to compute $P(\mathbf{x}, c_j)$ for each class c_j ?

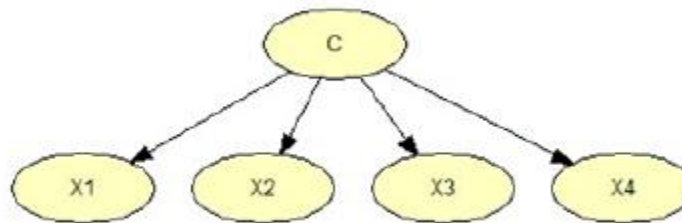
*Assuming complete data (all the attributes values for \mathbf{x} are known)
we do not need complicated inference algorithms
 \Rightarrow just calculate the joint probability distribution for each class*

Given a complete example $\mathbf{x} = (x_1, x_2, \dots, x_n)$

$$P(\mathbf{x}, c_j \mid S, \Theta_s) = \prod_{i=1}^n P(X_i = x_i \mid pa(X_i)). P(C = c_j \mid pa(C))$$

Naïve Bayes Bayesian Network

- A NB classifier can be viewed as a Bayesian network with a simple structure that has the class node as the parent node of all other attribute nodes.



- Priors $P(C)$ and conditionals $P(X_i | C)$ provide CPTs for the network.

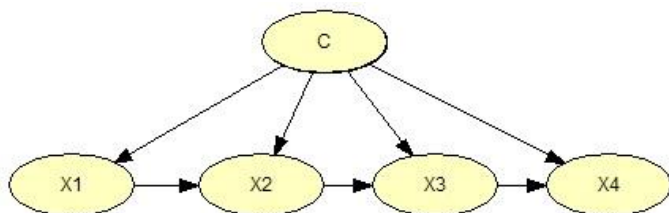
$$P(\mathbf{x}, c_j | S, \Theta_S) = \prod_{i=1}^n P(X_i = x_i | C = c_j) \cdot P(C = c_j)$$

because the structure of the Bayesian Network:

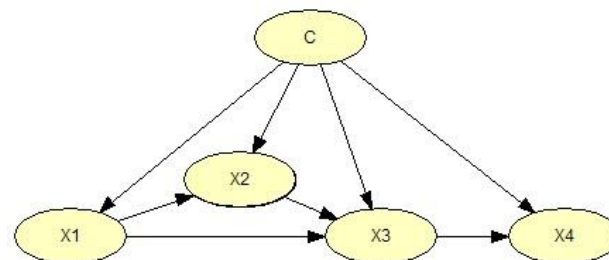
Bayesian Network Classifiers

Restricted vs. Unrestricted Approaches

- ▶ Restricted approach: contains the NB structure



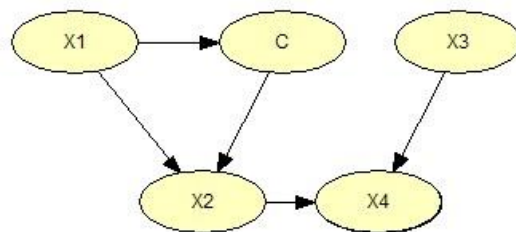
A **T**ree **A**ugmented **NB** (TAN)



A **BN** **A**ugmented **NB** (BAN)

(adding the best set of augmented arcs is an intractable problem)

- ▶ Unrestricted approach: the class node is treated as an ordinary node



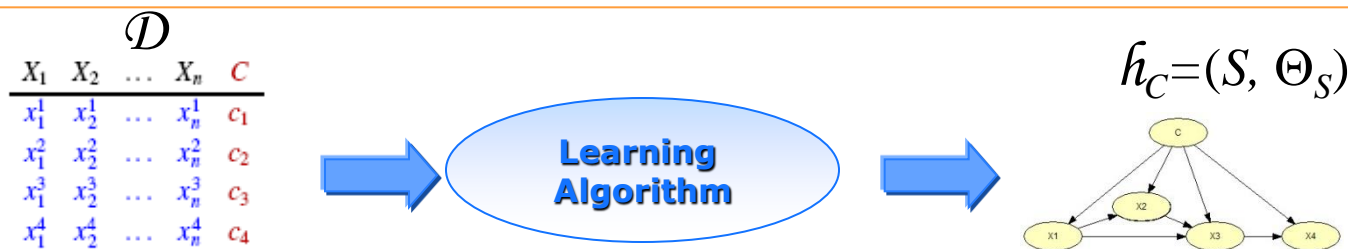
A General **BN**
(GBN) structure

Bayesian Network Classifiers

Learning Problem

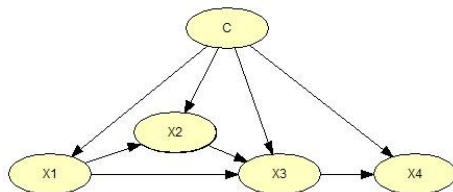
Given: a dataset \mathcal{D} of i.i.d labelled examples

Induce: the BNC that best fit the data in some sense



How to solve?

1. Choose a suitable class (class-model) of BNCs
 - For example a BAN can be chosen \Rightarrow the solution space is restricted
2. Choose a structure within this class-model that best fit the given data

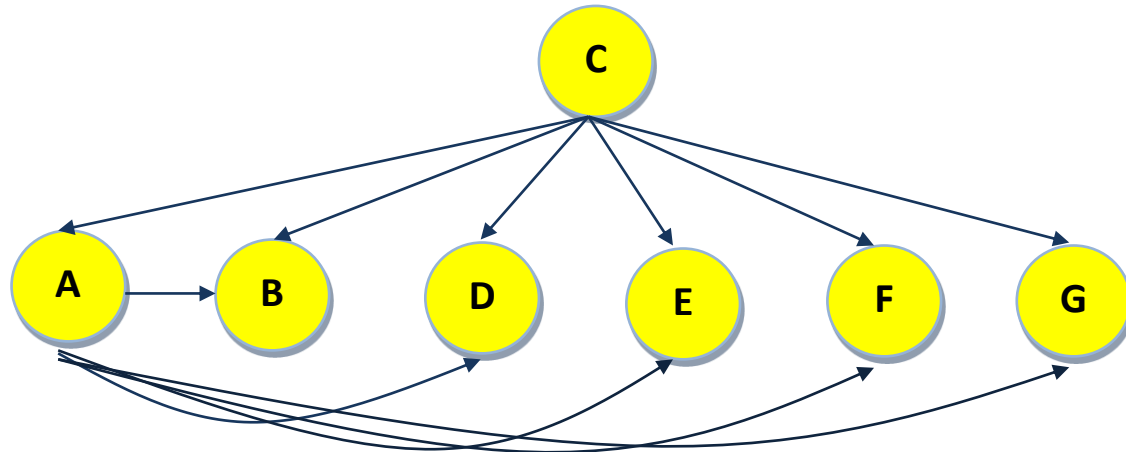


3. Estimate the parameters for the chosen structure from data

A Tree Augmented NB (TAN)

Friedman, Geiger and Goldszmith, 1997

- ▶ The attribute independence assumption is relaxed
⇒ less restrictive than Naive Bayes
- ▶ impose acceptable restriction: each attribute has as parents the class variable and at most one other attribute



polynomial time bound on constructing network
 $O((\# \text{ attributes})^2 * |\text{training set}|)$



TAN Learning Algorithm

(dependency analysis & search)

1. Compute the conditional mutual information, $I(X_i, X_j \mid C)$, between each pair of attributes, $i \neq j$
2. Build a complete undirected graph between all the attributes (excluding the class variable). Arc weight is the conditional mutual information between the vertices
3. Find maximum weight spanning tree over the graph
 1. Mark the two arcs with biggest weights and add to the tree
 2. Repeat until $n-1$ arc were added to the tree:
 - find the arc with biggest weight and add to the tree if not lead to a cycle
4. Transform the undirected tree to directed by picking a root in tree and making all arcs directed (to be outward from the root)
5. Construct a TAN model by adding a node labeled C and adding an arc from each attribute node to C



Mutual Information

Mutual Information: measures the mutual dependence of two variables. From the information theory point of view it measures the amount of information that can be obtained about one random variable by observing another.

$$I(X_i, X_j) = \sum_{i=1}^{r_x} \sum_{j=1}^{r_y} p(x_i, x_j) \log \frac{p(x_i, x_j)}{p(x_i) p(x_j)}$$

Conditional mutual information: measures the mutual information of two random variables conditioned on a third.

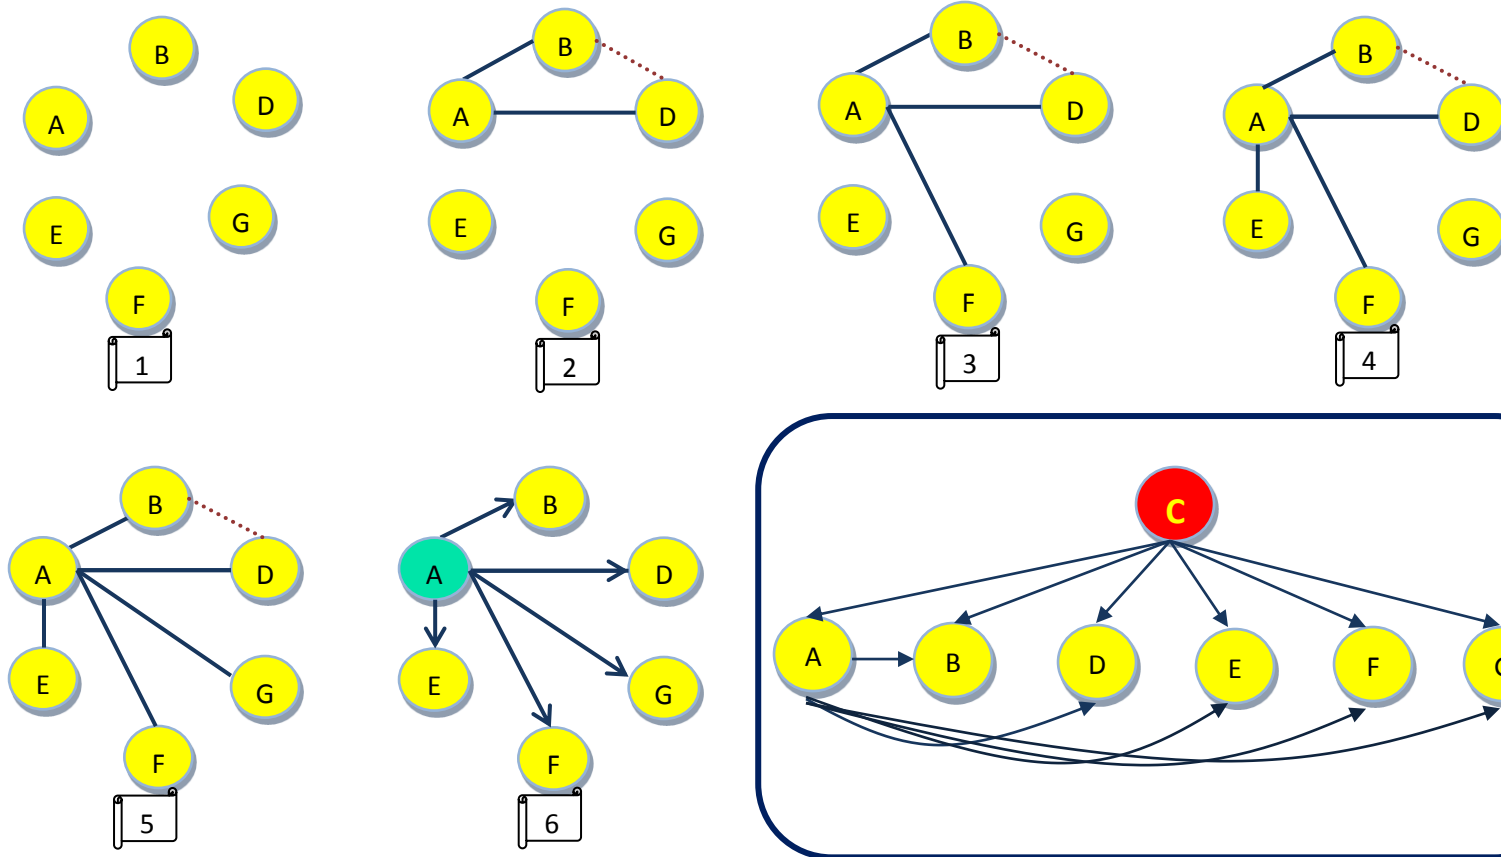
$$I(X_i, X_j | C) = \sum_c p(c) I(X_i, X_j | C = c) = \sum_{i=1}^{r_x} \sum_{j=1}^{r_y} \sum_{k=1}^{r_c} p(x_i, x_j, c_k) \log \frac{p(x_i, x_j | c_k)}{p(x_i | c_k) p(x_j | c_k)}$$

TAN Learning Algorithm

An example

6 attributes: A,B,D,E,F,G

$I(A,B|C) > I(A,D|C) > I(B,D|C) > I(A,F|C) > I(A,E|C) > I(A,G|C) > I(B,E|C) > I(B,F|C) > I(B,G|C) > I(D,E|C) > I(D,F|C) > I(D,G|C) > I(E,F|C) > I(E,G|C) > I(F,G|C)$

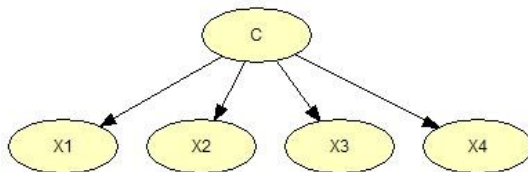


k-Dependence Bayesian Classifiers (k-DBC)

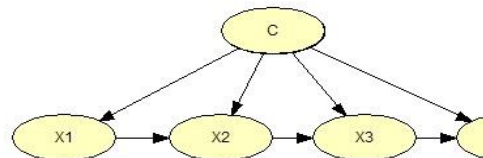
Sahami, M., 1996

k-DBC represents a unified framework for all the BNCs class-models containing the structure of the Naïve Bayes

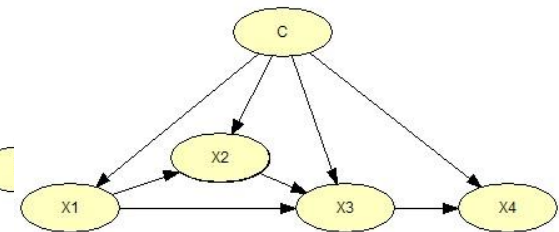
- Definition: A k-DBC is a Bayesian network which:
 - contains the structure of NB
 - allows each attribute X_i to have a maximum of k attribute nodes as parents



NB is a 0-DBC



TAN is a 1-DBC



This BAN is a 2-DBC

Model's complexity increases

By varying the value of k we can obtain classifiers that smoothly move along the spectrum of attribute dependences

We can control the complexity of k-DBC by choosing an appropriate k value



Sahami's Learning Algorithm

(dependency analysis & search)

- The algorithm can be viewed as a generalization of the TAN algorithm
- Mutual information is used as a measure of degree of attribute dependence
- Rationale:
 1. Starting with a k -DBC's structure S with a single class node C , the algorithm iteratively add $m = \min(|S|, k)$ parents to each new attribute added to S with largest dependence with the class C .
 2. The m parents for each new attribute are selected among those with higher degree of dependence given the class.
 3. The process finishes when all the attributes have been added to the structure S .

Sahami's Learning Algorithm

(dependency analysis & search)

Algorithm 5 Sahami's algorithm for learning k -DBCS

Require: A dataset \mathcal{D} of N labeled examples of $\langle \mathbf{X}, C \rangle$, the k value for the maximum allowable degree of attribute dependence

Ensure: A k -DBC

- 1: $V \leftarrow C$ {the set of nodes for the k -DBC}
 - 2: $A \leftarrow \emptyset$ {the set of arcs for the k -DBC}
 - 3: $\text{Temp} \leftarrow \emptyset$ {the used attribute list}
 - 4: **for all** attributes X_i and pair of attributes (X_i, X_j) such that $X_i \neq X_j$ **do**
 - 5: Compute $I(X_i, C)$ from data \mathcal{D}
 - 6: Compute $I(X_i, X_j | C)$ from data \mathcal{D}
 - 7: **repeat**
 - 8: Select X_{max} such that $X_{max} = \arg \max_{X_i \notin \text{Temp}} I(X_i, C)$
 - 9: Add the node X_{max} to V
 - 10: Add the arc (C, X_{max}) to A
 - 11: Add $m = \min(|\text{Temp}|, k)$ arcs to A from m distinct attributes $X_j \in \text{Temp}$ with the highest value of $I(X_i, X_j | C)$
 - 12: Add the attribute X_{max} to Temp
 - 13: **until** Temp includes all the attributes $X_i \in \mathbf{X}$
 - 14: Compose S such that $S = (V, A)$ {the k -DBC structure}
 - 15: Estimate the parameters Θ_S given S from data \mathcal{D}
 - 16: **return** $k\text{-DBC} = (S, \Theta_S)$
-



Learning BNCs augmented NB Hill Climbing using only arc addition

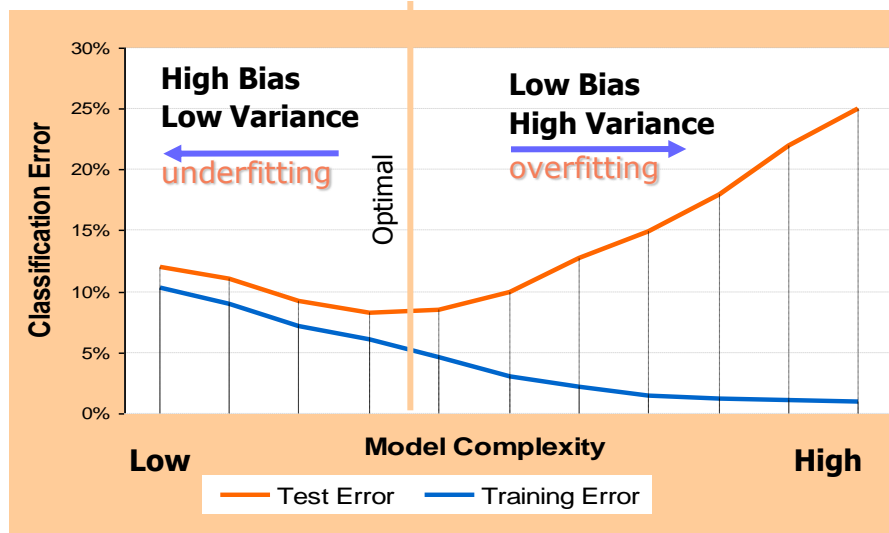
(scoring & search)

```
procedure Learn_BNC(data, maxNrOfParents, score){  
  Init:   BNC ← Learn_NaiveBayes(data)  
          bestScore ← calcScore(BNC, score)  
  
  repeat  // go do the search  
    bestArc ← findBestArcToAdd(BNC, data, maxNrOfParents)  
    newScore ← calcScore(BNC, score);  
    if (bestArc != null) && (newScore > bestScore)  
      BNC ← AddArc (BNC, bestArc)  
      bestScore ← newScore  
    else  
      return(BNC)  
}
```

Choosing the Class-Model

Learning algorithms should be able to select a model with the appropriate complexity for the available data

Behavior of test error and training error
varying model complexity



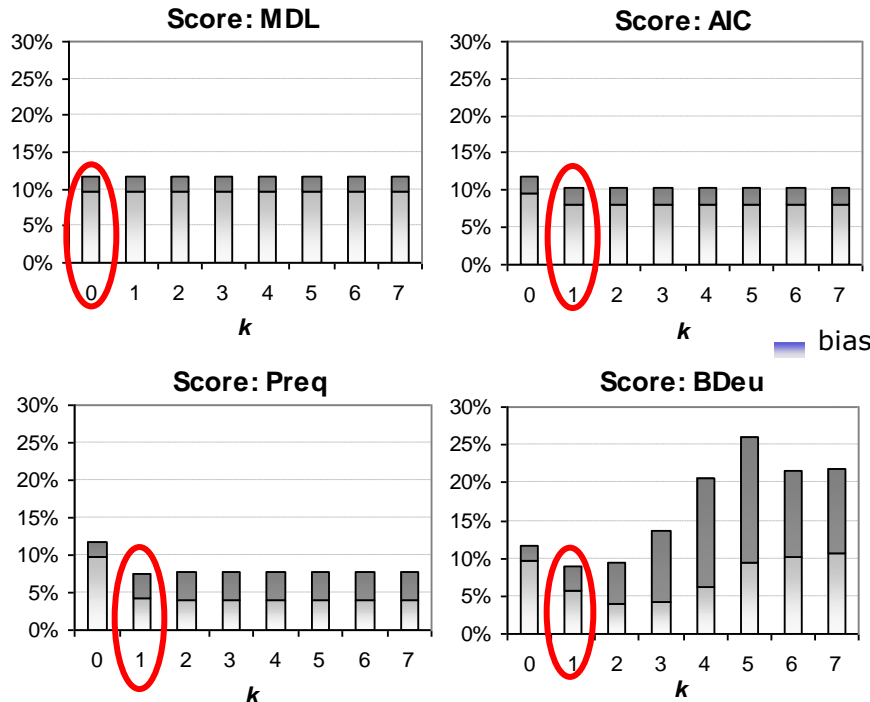
there is an optimal model that gives minimum test error

- Classifier too simple (e.g. NB)
 - underfit the data, too much bias
- Classifier too complex
 - overfit the data, too much variance
- Controlling the BNCs complexity
 - BNC class-models: NB, TAN, BAN, etc., differ by the number of parents allowed for attribute
 - choosing of the appropriate class-model of BNCs \Rightarrow **not trivial**: depend on the chosen score and available data

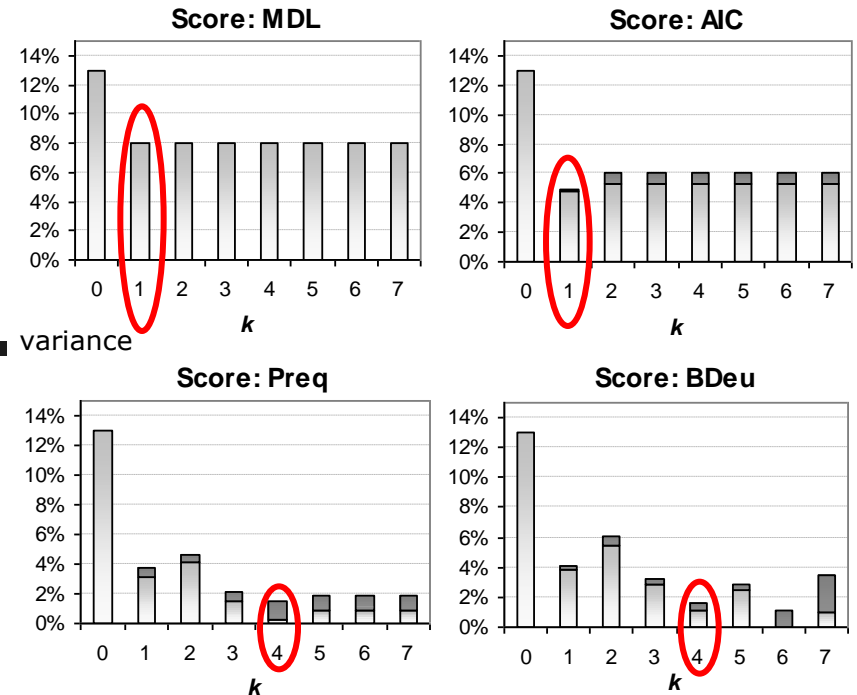
What k to choose?

Varying k , the score and the training set size can have different effects on bias and variance and consequently in the test error

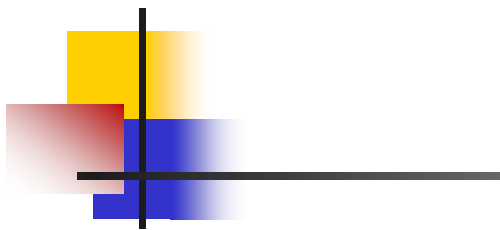
1000 training examples



12500 training examples



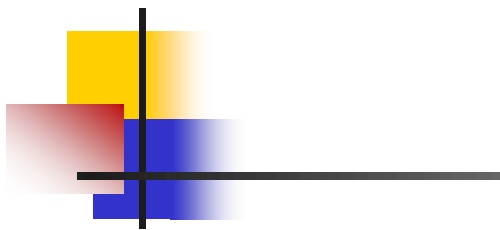
If the data increases, it makes sense to **gradually increase** the k value to adjust the complexity of the class-model, and hence, the complexity of BNCs to the available data



k-DBCs learned from the Nursery dataset using the hill-climbing learning algorithm with five unsupervised scores:

- MLC
- MDL
- AIC
- Bayes
- BDe





k-DBC's learned from the Nursery dataset using the hill-climbing learning algorithm with three supervised scores:

- LOO-CV
- k-Fold-CV
- Cumulative



Learning BNC with RapidMiner

The screenshot displays the RapidMiner GUI with the following components:

- Top Bar:** Menu (File, Edit, Process, Tools, View, Help) and a toolbar with icons for file operations and execution.
- Left Panel:** A tree view under 'wBayes' showing a hierarchy: Modeling (10) > Classification and Regression (10) > Weka (10) > Bayes (8). The 'W-BayesNet' operator is selected and highlighted.
- Central Canvas:** Titled 'Main Process', it shows a workflow diagram with a 'Retrieve' operator connected to a 'W-BayesNet' operator. The 'W-BayesNet' operator has ports labeled 'inp', 'out', 'tra', 'mod', and 'exa'.
- Right Panel:** The 'Parameters' tab for 'W-BayesNet'. It lists several parameters with input fields:
 - D: -Q
 - B: (empty)
 - Q: weka.classifiers.bayes
 - E: weka.classifiers.bayes
- Bottom Panel:** Includes a 'Remote Processes' section, a 'Problems' tab showing 'No problems found', and a 'Log' tab.
- Bottom Right:** A 'Help' pane titled 'Weka:W-BayesNet' containing a 'Synopsis' and a 'Description' of the operator.

Building a TAN classifier using the Weka class W-BayesNet

Parameter Settings:

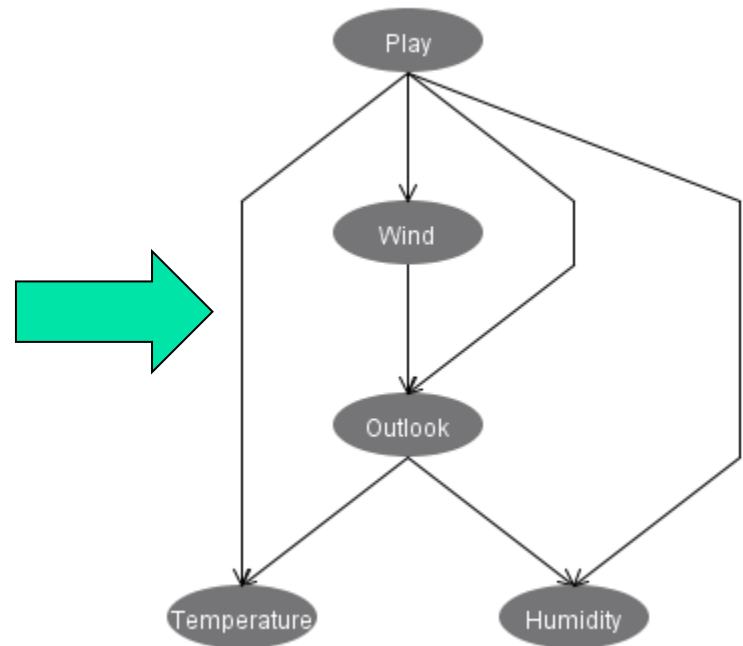
Q = weka.classifiers.bayes.net.search.local.TAN (structure learning algorithm)

E = weka.classifiers.bayes.net.estimate.SimpleEstimator (parameter learning algorithm)

Golf Dataset

ExampleSet (14 examples, 1 special attribute, 4 regular attributes)

Row No.	Play	Outlook	Temperature	Humidity	Wind
1	no	sunny	85	85	false
2	no	sunny	80	90	true
3	yes	overcast	83	78	false
4	yes	rain	70	96	false
5	yes	rain	68	80	false
6	no	rain	65	70	true
7	yes	overcast	64	65	true
8	no	sunny	72	95	false
9	yes	sunny	69	70	false
10	yes	rain	75	80	false
11	yes	sunny	75	70	true
12	yes	overcast	72	90	true
13	yes	overcast	81	75	false
14	no	rain	71	80	true



Learning BNCs using the Weka class `weka/classifiers/bayes/BayesNet`

Several learning algorithms for BNCs are implemented in Weka
(see full documentation at <http://www.cs.waikato.ac.nz/~remco/weka.bn.pdf>)

The screenshot displays the Weka graphical user interface. On the left, the 'Main Process' window shows a workflow starting with an 'inp' node connected to a 'Retrieve' node, which then connects to a 'W-BayesNet' node. The 'W-BayesNet' node has four sub-nodes: 'tra', 'mod', 'exa', and 'res'. The 'res' node is connected to a 'res' output node. On the right, the 'Parameters' window for 'W-BayesNet' is open, showing the following settings:

Parameter	Value
D	-Q
B	
Q	assifiers.bayes.net.search.local.TAN
E	bayes.net.estimate.SimpleEstimator

At the bottom, a 'Help' window is open, showing the 'Weka:W-BayesNet' synopsis:

Synopsis
Bayes Network learning using various search algorithms and quality measures.



References

1. Buntine, W. (1996) A guide to the literature on learning probabilistic networks from data, IEEE Transactions on Knowledge and Data Engineering 8, 195-210
2. Bouckaert, R. (2007) Bayesian Network Classifiers in Weka, Technical Report
3. Castillo, G. (2006) Adaptive Learning Algorithms for Bayesian Network Classifiers, PhD Thesis, Chapter 3.
4. Cheng, J. and Greiner, R. (1999) Comparing Bayesian Network Classifiers, Proceedings of UAI
5. Chickering, D.M. (2002) Learning equivalence classes of Bayesian-network structures. *Journal of Machine Learning Research*
6. Friedman, N., Geiger, D. and Goldszmidt M. (1997) Bayesian network classifiers, Machine Learning 29, 139-164.
7. Heckerman D. (1995) A tutorial on learning with Bayesian Networks. Technical Report. Microsoft Research
8. Mitchell T. (1997) Machine Learning, Chapter 6. Bayesian Learning, Tom Mitchell, McGraw Hill

Presentation Slides

- ❖ [Learning Bayesian Networks from Data](#) by Nir Friedman and Daphne Koller
 - ❖ [Powerpoint Presentation](#)
 - ❖ [Document in PDF \(1.72MB\)](#)
- ❖ [Learning Bayesian Networks from Data by](#) Nir Friedman and Moises Goldszmidt
 - ❖ [Powerpoint Presentation Show \(1.4MB\), PDF \(15MB\) \(6 x pages\)](#)
- ❖ [Learning Bayesian Networks](#) by Andrew Moore
- ❖ [Aprendizagem Bayesiana](#) by João Gama, Universidade do Porto
- ❖ [Clasificadores Bayesianos](#) by Abdelmalik Moujahid, Iñaki Inza e Pedro Larrañaga. Universidad del País Vasco