


Branch: master ▾

Find file Copy path

pybni / examples / 010-2018-06-15-marks.ipynb

 **cs224** 010-2018-06-15-marks.ipynb: current version is also printing the bnle...

99d5772 on Jun 22, 2018

1 contributor



Raw Blame History



3055 lines (3055 sloc) 114 KB

In [1]:

```
%load_ext watermark
%watermark -a 'Christian Schuhegger' -u -d -v -p n
umpy,xarray,scipy,pandas,sklearn,matplotlib,seabor
n,qgrid,rpy2,libpgm,pgmpy,networkx,graphviz,pybml,
pytest
```

Christian Schuhegger
last updated: 2018-06-22

CPython 3.6.4
IPython 6.2.1

numpy 1.14.2
xarray 0.10.3
scipy 1.0.1
pandas 0.22.0
sklearn 0.19.1
matplotlib 2.2.2
seaborn 0.8.1
qgrid 1.0.2
rpy2 2.9.1
libpgm n
pgmpy n
networkx 2.1
graphviz 0.8.3
pybml n
pytest 3.5.0

Access other versions via nbviewer:

<https://nbviewer.jupyter.org/github/cs224/pybml/blob/4bd08e6e48194dcddeadf0f202b910e3e224753b/examples/010-2018-06-15-marks.ipynb>

(<https://nbviewer.jupyter.org/github/cs224/pybml/blob/4bd08e6e48194dcddeadf0f202b910e3e224753b/examples/010-2018-06-15-marks.ipynb>)

In [2]:

```
%matplotlib inline
import numpy as np, pandas as pd, xarray as xr, ma
tplotlib.pyplot as plt, seaborn as sns
import sklearn, sklearn.pipeline
import networkx as nx, graphviz, networkx.algorith
ms.dag
import random

pd.set_option('display.max_columns', 500)
pd.set_option('display.width', 1000)
# pd.set_option('display.float_format', lambda x:
'%.2f' % x)
np.set_printoptions(edgeitems=10)
np.set_printoptions(suppress=True)
np.core.arrayprint._line_width = 180

sns.set()
```

In [3]:

```

from IPython.display import display, HTML

from IPython.display import display_html
def display_side_by_side(*args):
    html_str=''
    for df in args:
        if type(df) == np.ndarray:
            df = pd.DataFrame(df)
            html_str+=df.to_html()
        html_str = html_str.replace('table','table style="display:inline"')
        # print(html_str)
        display_html(html_str,raw=True)

CSS = """
.output {
    flex-direction: row;
}
"""

def display_graphs_side_by_side(*args):
    html_str='<table><tr>'
    for g in args:
        html_str += '<td>'
        html_str += g._repr_svg_()
        html_str += '</td>'
    html_str += '</tr></table>'
    display_html(html_str,raw=True)

display(HTML("<style>.container { width:70% !important; }</style>"))

```

In [5]:

```
%load_ext rpy2.ipython
```

The rpy2.ipython extension is already loaded. To reload it, use:

```
%reload_ext rpy2.ipython
```

In [6]:

```

%load_ext autoreload
%autoreload 1
%aimport pybml.bn

```

In [7]:

```

import locale
locale.setlocale(locale.LC_ALL, 'C')

import rpy2, rpy2.rinterface, rpy2.robjects, rpy2.
robjects.packages, rpy2.robjects.lib, rpy2.robject

```


```
s.lib.grid, \
    rpy2.robjjects.lib.ggplot2, rpy2.robjjects.panda
s2ri, rpy2.interactive.process_revents, \
    rpy2.interactive, rpy2.robjjects.lib.grdevices
# rpy2.interactive.process_revents.start()
rpy2.robjjects.pandas2ri.activate()
```

In [9]:

```
rpackageversionfn = rpy2.robjjects.r('packageVersio
n')
print(rpackageversionfn("bnlearn")[0])
```

```
[1]          4          4 20180620
```

learning.test

Before we look at the marks data-set let's look first at a test network provided in the bnlearn package: [networks](http://www.bnlearn.com/documentation/networks/) (<http://www.bnlearn.com/documentation/networks/>) 
 src='http://www.bnlearn.com/documentation/networks/learning.test.png' (<http://www.bnlearn.com/documentation/networks/learning.test.png>)' width=400>

In [10]:

```
%%R -o rdf_lt
data(learning.test)
rdf_lt = learning.test
```

Converting the R data.frame into a python pd.DataFrame and converting to CategoricalDtype

After loading the data-set we need to convert it so that all variables are of type CategoricalDtype: see the pandas documentation [about Categorical Data](https://pandas.pydata.org/pandas-docs/stable/categorical.html) (<https://pandas.pydata.org/pandas-docs/stable/categorical.html>) for more details.

In [11]:

```
#df_lt = rpy2.robjjects.pandas2ri.r2py(rdf_lt)
df_lt = rdf_lt

ct1 = pd.api.types.CategoricalDtype(['a', 'b', 'c'], ordered=True)
ct2 = pd.api.types.CategoricalDtype(['a', 'b'], ordered=True)

for c in 'ABCDE':
    df_lt[c] = df_lt[c].astype(ct1)
```

```
df_lt['F'] = df_lt['F'].astype(ctz)
```

```
df_lt.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 5000 entries, 1 to 5000
```

```
Data columns (total 6 columns):
```

```
A    5000 non-null category
```

```
B    5000 non-null category
```

```
C    5000 non-null category
```

```
D    5000 non-null category
```

```
E    5000 non-null category
```

```
F    5000 non-null category
```

```
dtypes: category(6)
```


```
memory usage: 69.0+ KB
```

ConstraintBasedNetFromDataDiscreteBayesNetwork

In [12]:

```
cbnet = pybnl.bn.ConstraintBasedNetFromDataDiscreteBayesNetwork(df_lt)
cbnet.fit()
#display_side_by_side(cbnet.structure().dot(), cbnet.structure().cpdag().dot())
cbnet.structure().cpdag().dot()
```

Out[12]:





ScoreBasedNetFromDataDiscreteBayesNetwork

In [13]:

```
sbnet = pybnl.bn.ScoreBasedNetFromDataDiscreteBayesNetwork(df_lt)
sbnet.fit()
sbnet.structure().cpdag().dot()
```

Out[13]:





HybridScoreAndConstainedBasedNetFromDataDiscreteBayes


In [14]:

```

hnet1 = pybnl.bn.HybridScoreAndConstainedBasedNetFromDataDiscreteBayesNetwork(df_lt)
hnet1.fit()
hnet1.structure().cpdag().dot()

```

Out[14]:




In [15]:

```

hnet2 = pybnl.bn.HybridScoreAndConstainedBasedNetFromDataDiscreteBayesNetwork(df_lt, algorithm='rxmax2_sihitonpc_tabu')
hnet2.fit()
hnet2.structure().cpdag().dot()

```

Out[15]:



In [16]:

```

hnet2.structure().cpdag().vstructs()

```

Out[16]:

```

[[A, B], [A, D], [B, E], [C, D], [F, E]]

```

	X	Z	Y
0	A	D	C
1	B	E	F

marks

Let's take the detour of loading the R data set, writing it to CSV and then loading the CSV via pandas from python. Like that we're sure we have a typical starting position in a python data workflow.

In [17]:

```
%%R -o marks
library(bnlearn)
data(marks)
write.csv(marks, file = "marks.csv")
```

In [18]:

```
pd_marks = pd.read_csv('marks.csv', index_col=0).a
stype(np.float64)
pd_marks.head()
```

Out[18]:

	MECH	VECT	ALG	ANL	STAT
1	77.0	82.0	67.0	67.0	81.0
2	63.0	78.0	80.0	70.0	81.0
3	75.0	73.0	71.0	66.0	81.0
4	55.0	72.0	63.0	70.0	68.0
5	63.0	63.0	65.0	70.0	63.0

In [19]:

```
hbt = pybml.bn.HarteminkBinTransformer(3,ibreaks=1
8)
dmarks = hbt.fit_transform(X=pd_marks)
```

In [20]:

```
# dmarks = pybml.bn.discretize(pd_marks)
dmarks.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 88 entries, 0 to 87
Data columns (total 5 columns):
MECH      88 non-null category
VECT      88 non-null category
```

```

VECT      88 non-null category
ALG       88 non-null category
ANL       88 non-null category
STAT      88 non-null category
dtypes: category(5)
memory usage: 1.0 KB

```

In [21]:

```
dmarks['MECH'].dtype
```

Out[21]:

```
CategoricalDtype(categories=['(0,35.8]', '(35.8,49]', '(49,77]'], ordered=False)
```

In [22]:

```
dmarks.head()
```

Out[22]:

	MECH	VECT	ALG	ANL	STAT
0	(49,77]	(60.7,82]	(58.7,80]	(62.3,70]	(53.7,81]
1	(49,77]	(60.7,82]	(58.7,80]	(62.3,70]	(53.7,81]
2	(49,77]	(60.7,82]	(58.7,80]	(62.3,70]	(53.7,81]
3	(49,77]	(60.7,82]	(58.7,80]	(62.3,70]	(53.7,81]
4	(49,77]	(60.7,82]	(58.7,80]	(62.3,70]	(53.7,81]

Let's also create immediately a marks data-frame that include one additional latent variable that we will need later:

In [23]:

```

ldmarks = dmarks.copy()
pybml.bn.augment_df_with_latent_variable(ldmarks,
'LAT', 3)
print(pybml.bn.levels_of_latent_variable(ldmarks,
'LAT'))
ldmarks.head()

```

```
['1000', '1001', '1002']
```

Out[23]:

	MECH	VECT	ALG	ANL	STAT	LAT
0	(49,77]	(60.7,82]	(58.7,80]	(62.3,70]	(53.7,81]	NaN
1	(49,77]	(60.7,82]	(58.7,80]	(62.3,70]	(53.7,81]	NaN
2	(49,77]	(60.7,82]	(58.7,80]	(62.3,70]	(53.7,81]	NaN
3	(49,77]	(60.7,82]	(58.7,80]	(62.3,70]	(53.7,81]	NaN
4	(49,77]	(60.7,82]	(58.7,80]	(62.3,70]	(53.7,81]	NaN

NetAndDataDiscreteBayesNetwork

Create network by hand

In [24]:

```
dg = nx.DiGraph()
# G.add_node(1)
dg.add_nodes_from(list(marks.columns))
dg.add_edges_from([
    ['STAT', 'ANL'],
    ['STAT', 'ALG'],
    ['ANL', 'ALG'],
    ['ALG', 'MECH'],
    ['ALG', 'VECT'],
    ['VECT', 'MECH'],
])
```

In [25]:

```
list(nx.connected_components(dg.to_undirected()))
```

Out[25]:

```
[{'ALG', 'ANL', 'MECH', 'STAT', 'VECT'}]
```

In [26]:

```
ns = pybml.bn.digraph2netstruct(dg)
```

You can display the graph either in the 'dot' format or in the default more compact 'fdp' format

In [27]:


```
type(ns.dot())
```


Out[27]:

```
graphviz.dot.Digraph
```

In [28]:

```
display_graphs_side_by_side(ns.dot(engine='dot'),
ns.dot())
```





Fit the network

In [29]:

```
net_dmarks = pybnl.bn.NetAndDataDiscreteBayesNetwork(
    dmarks, rnet=ns.rnet)
net_dmarks.fit()
net_dmarks_net_and_data = net_dmarks
print(net_dmarks.rfit)
```

Bayesian network parameters

Parameters of node ALG (multinomial distribution)

Conditional probability table:

, , STAT = (9,25.5]

	ANL	
ALG	(9,39.3]	(39.3,62.3]
(15,48.8]	1.0000000	0.5000000
(48.8,58.7]	0.0000000	0.5000000
(58.7,80]	0.0000000	0.0000000

, , STAT = (25.5,53.7]

	ANL	
ALG	(9,39.3]	(39.3,62.3]
(15,48.8]	0.7692308	0.2250000
(48.8,58.7]	0.2307692	0.6250000
(58.7,80]	0.0000000	0.1500000

, , STAT = (53.7,81]

	ANL		
ALG	(9,39.3]	(39.3,62.3]	(62.3,70]
(15,48.8]	1.0000000	0.1111111	0.0000000
(48.8,58.7]	0.0000000	0.4444444	0.0000000
(58.7,80]	0.0000000	0.4444444	1.0000000

Parameters of node ANL (multinomial distribution)

Conditional probability table:

	STAT		
ANL	(9,25.5]	(25.5,53.7]	(53.7,81]
(9,39.3]	0.7333333	0.2452830	0.0500000
(39.3,62.3]	0.2666667	0.7547170	0.4500000
(62.3,70]	0.0000000	0.0000000	0.5000000

Parameters of node MECH (multinomial distribution)

Conditional probability table:

, , VECT = (9,42.2]

	ALG		
MECH	(15,48.8]	(48.8,58.7]	(58.7,80]
(0,35.8]	0.7777778	0.5000000	0.0000000
(35.8,49]	0.2222222	0.3333333	1.0000000
(49,77]	0.0000000	0.1666667	0.0000000

, , VECT = (42.2,60.7]

	ALG		
MECH	(15,48.8]	(48.8,58.7]	(58.7,80]
(0,35.8]	0.3333333	0.3333333	0.57142857
(35.8,49]	0.4000000	0.57142857	0.14285714
(49,77]	0.2666667	0.09523810	0.28571429

, , VECT = (60.7,82]

	ALG		
MECH	(15,48.8]	(48.8,58.7]	(58.7,80]
(0,35.8]	0.0000000	0.14285714	0.0000000
(35.8,49]	1.0000000	0.57142857	0.08333333
(49,77]	0.0000000	0.28571429	0.91666667

Parameters of node STAT (multinomial distribution)

Conditional probability table:

(9,25.5]	(25.5,53.7]	(53.7,81]
0.1704545	0.6022727	0.2272727

Parameters of node VECT (multinomial distribution)

Conditional probability table:

	ALG		
VECT	(15,48.8]	(48.8,58.7]	(58.7,80]
(9,42.2]	0.52941176	0.17647059	0.05000000


```
(42.2,60.7] 0.44117647 0.61764706 0.35000000  
(60.7,82] 0.02941176 0.20588235 0.60000000
```

ConstraintBasedNetFromDataDiscreteBayesNetwork

In [30]:

```
net_dmarks = pybnl.bn.ConstraintBasedNetFromDataDiscreteBayesNetwork(dmarks)  
net_dmarks.fit()  
net_dmarks_cb = net_dmarks  
net_dmarks.structure().cpdag().dot()
```

Out[30]:





ScoreBasedNetFromDataDiscreteBayesNetwork

In [31]:

```
net_dmarks = pybnl.bn.ScoreBasedNetFromDataDiscreteBayesNetwork(dmarks)  
net_dmarks.fit()  
net_dmarks_sb = net_dmarks  
net_dmarks.structure().cpdag().dot()
```

Out[31]:





HybridScoreAndConstainedBasedNetFromDataDiscreteBayes


mmhc

In [32]:

```

net_dmarks = pybnl.bn.HybridScoreAndConstainedBase
dNetFromDataDiscreteBayesNetwork(dmarks)
net_dmarks.fit()
net_dmarks_hyb_mmhc = net_dmarks
net_dmarks.structure().cpdag().dot()
  
```

Out[32]:



rxmax2_sihitonpc_tabu


In [33]:

```

net_dmarks = pybnl.bn.HybridScoreAndConstainedBase
dNetFromDataDiscreteBayesNetwork(dmarks, algorithm
='rxmax2_sihitonpc_tabu')
net_dmarks.fit()
net_dmarks_hyb_rsmax = net_dmarks
net_dmarks.structure().cpdag().dot()
  
```

Out[33]:





Comparison

In [34]:

```
def score(type='bic'):
    cdf = pd.DataFrame(columns=['algorithm', 'score'])
    cdf.loc[len(cdf)] = ['ConstraintBased', net_dmarks_cb.score(type=type)]
    cdf.loc[len(cdf)] = ['ScoreBased', net_dmarks_sb.score(type=type)]
    cdf.loc[len(cdf)] = ['mmhc', net_dmarks_hyb_mmhc.score(type=type)]
    cdf.loc[len(cdf)] = ['rsmax', net_dmarks_hyb_rsmax.score(type=type)]
    return cdf

display_side_by_side(score(), score(type='loglik'))
```

	algorithm	score
0	ConstraintBased	-411.497387
1	ScoreBased	-411.497387
2	mmhc	-452.330225
3	rsmax	-434.906936

	algorithm	score
0	ConstraintBased	-353.292008
1	ScoreBased	-353.292008
2	mmhc	-412.034194
3	rsmax	-394.610904

In [35]:

```
net_dmarks_cb.structure().bf(net_dmarks_hyb_mmhc.structure(), dmarks)
```

Out[35]:

45.10970279456899

In [36]:

```
net_dmarks_cb.bf(net_dmarks_hyb_mmhc, dmarks)
```

Out[36]:

45.10970279456899

In [37]:

```
net_dmarks_cb.bf(net_dmarks_hyb_mmhc)
```

Out[37]:

45.10970279456899

sklearn pipeline

In [38]:

```
pl = sklearn.pipeline.Pipeline(  
    steps=[  
        ('hartemink', pybml.bn.HarteminkBinTransfo  
rmer(3,ibreaks=18)),  
        ('score', pybml.bn.ScoreBasedNetFromDataDi  
screteBayesNetwork()),  
    ]  
)  
pl.set_params(hartemink__breaks=3).fit(pd_marks, N  
one)
```

Out[38]:

```
Pipeline(memory=None,  
    steps=[('hartemink', HarteminkBinTransformer  
(breaks=3, ibreaks=18)), ('score', ScoreBasedNetFr  
omDataDiscreteBayesNetwork(algorithm=None, ldf=Non  
e,  
                                whitelist=None))])
```

In [39]:

```
pl.score(dmarks)
```


Out[39]:


-411.49738682714093

In [40]:

```
pl.steps[-1][1].structure().dot()
```

Out[40]:





In [41]:

```
pl.set_params(hartemink__breaks=5).fit(pd_marks, N
one)
pl.score(dmarks)
```


Out[41]:

-682.3887051571114

In [42]:

```
pl.steps[-1][1].structure().dot()
```

Out[42]:



In [43]:

```
pl.set_params(hartemink__breaks=7).fit(pd_marks, N
one)
pl.score(dmarks)
```

Out[43]:


-876.9315865210814

In [44]:

```
pl.steps[-1][1].structure().dot()
```

Out[44]:





StructuralEMNetFromDataDiscreteBavesNetwork